

Software Foundations of Security & Privacy

15315 Spring 2017


Lecture 4:

Enforceable Security Policies

Matt Fredrikson
mfredrik@cs.cmu.edu

January 26, 2017

Securing Extensible Systems



Key Questions

What security policies can we enforce?

- ▶ Topic of today's lecture

What mechanisms can we use?

Key Questions

What security policies can we enforce?

- ▶ Topic of today's lecture

What mechanisms can we use?

- ▶ Type checking
- ▶ Static verification
- ▶ Program rewriting
- ▶ **Runtime enforcement**

Our focus: policies enforceable by **execution monitors**

Our focus: policies enforceable by **execution monitors**

Execution Monitor (EM)

An execution monitor is a coroutine that executes in parallel with a **target** program or system.

- ▶ Monitor steps of a *single* execution
- ▶ Compare observed behavior against a policy
- ▶ Terminate the program when policy is violated

Execution monitor examples



Execution monitor examples

- ▶ Filesystem access control
- ▶ Firewall
- ▶ Stack inspection
- ▶ Dynamic bounds checking
- ▶ Malware detectors
- ▶ Chrome's *Content Security Policies* (CSP)
- ▶ ...

What's *not* EM?

**anything that uses more information than what's available
from a single execution of the program/system**

What's *not* EM?

**anything that uses more information than what's available
from a single execution of the program/system**

In particular, this excludes:

- ▶ Information about *future* steps
- ▶ Alternative *hypothetical* executions
- ▶ *All possible* executions

What's *not* EM?

anything that uses more information than what's available from a single execution of the program/system

In particular, this excludes:

- ▶ Information about *future* steps
- ▶ Alternative *hypothetical* executions
- ▶ *All possible* executions

As we'll see this excludes some important properties

What's *not* EM?

anything that uses more information than what's available from a single execution of the program/system

In particular, this excludes:

- ▶ Information about *future* steps
- ▶ Alternative *hypothetical* executions
- ▶ *All possible* executions

As we'll see this excludes some important properties

Later, we'll talk about techniques that aren't limited in this way

- ▶ Verifying compilers, type systems
- ▶ Anything classified as “static analysis”

Formalizing execution

A target S is characterized by:

- ▶ A set of **atomic actions** A
- ▶ A set of sequences Σ_S of elements from A

Formalizing execution

A target S is characterized by:

- ▶ A set of **atomic actions** A
- ▶ A set of sequences Σ_S of elements from A

Sequences in Σ_S can be finite or infinite

Formalizing execution

A target S is characterized by:

- ▶ A set of **atomic actions** A
- ▶ A set of sequences Σ_S of elements from A

Sequences in Σ_S can be finite or infinite

What might A look like?

Formalizing execution

A target S is characterized by:

- ▶ A set of **atomic actions** A
- ▶ A set of sequences Σ_S of elements from A

Sequences in Σ_S can be finite or infinite

What might A look like?

- ▶ Set of program states: mappings from *variables* to *values*

Formalizing execution

A target S is characterized by:

- ▶ A set of **atomic actions** A
- ▶ A set of sequences Σ_S of elements from A

Sequences in Σ_S can be finite or infinite

What might A look like?

- ▶ Set of program states: mappings from *variables* to *values*
- ▶ Set of all system calls: open, send, ...

Formalizing execution

A target S is characterized by:

- ▶ A set of **atomic actions** A
- ▶ A set of sequences Σ_S of elements from A

Sequences in Σ_S can be finite or infinite

What might A look like?

- ▶ Set of program states: mappings from *variables* to *values*
- ▶ Set of all system calls: `open`, `send`, ...
- ▶ Set of primitive commands in server scripting language

Example

How do we model the following program?

```
while( $x < y$ ) {  
   $x := x + 1$ ;  
}
```

Example

How do we model the following program?

```
while( $x < y$ ) {  
   $x := x + 1$ ;  
}
```

Atomic actions $A = \{\text{mappings from } x, y \text{ to } \mathbb{Z}\}$

Example

How do we model the following program?

while($x < y$) {
 $x := x + 1$;
}

Atomic actions $A = \{\text{mappings from } x, y \text{ to } \mathbb{Z}\}$

$$\Sigma_S = \left\{ \begin{array}{l} \dots \\ [(x \mapsto -1, y \mapsto 0), (x \mapsto 0, y \mapsto 0)] \\ [(x \mapsto 0, y \mapsto 0)] \\ [(x \mapsto 0, y \mapsto 1), (x \mapsto 1, y \mapsto 1)] \\ [(x \mapsto 0, y \mapsto 2), (x \mapsto 1, y \mapsto 2), (x \mapsto 2, y \mapsto 2)] \\ \dots \end{array} \right\}$$

Example

How might we model the following program?

```
while(read(&buf, &len, fp)) {  
    if(buf[0] == 255)  
        send(sock, buf, len);  
    printf("%s", buf);  
}
```

Example

How might we model the following program?

```
while(read(&buf, &len, fp)) {  
    if(buf[0] == 255)  
        send(sock, buf, len);  
    printf("%s", buf);  
}
```

$$A = \{\text{read}, \text{send}\}$$

Example

How might we model the following program?

```
while(read(&buf, &len, fp)) {  
    if(buf[0] == 255)  
        send(sock, buf, len);  
    printf("%s", buf);  
}
```

$A = \{\text{read}, \text{send}\}$

$$\Sigma_S = \left\{ \begin{array}{l} [\text{read}, \text{read}, \dots] \\ [\text{read}, \text{send}, \dots] \\ [\text{read}, \text{send}, \text{read}, \dots] \\ \dots \end{array} \right\}$$

Formalizing policies

Let Ψ denote the universe of all possible executions in \mathcal{A}

- ▶ Note: Ψ is not the same as Σ_S
- ▶ It contains executions that may not be possible in S
- ▶ In particular, $\Sigma_S \subseteq \Psi$

Formalizing policies

Let Ψ denote the universe of all possible executions in A

- ▶ Note: Ψ is not the same as Σ_S
- ▶ It contains executions that may not be possible in S
- ▶ In particular, $\Sigma_S \subseteq \Psi$

Policy

A **policy** P is a predicate on *sets* of executions. In other words,

$$P \subseteq 2^\Psi$$

A target S satisfies P if and only if $\Sigma_S \in P$.

Example

Suppose that we want a simple policy

```
while(read(&buf, &len, fp)) {  
    if(buf[0] == 255)  
        send(sock, buf, len);  
    printf("%s", buf);  
}
```

“No send after read”

Example

Suppose that we want a simple policy

```
while(read(&buf, &len, fp)) {  
    if(buf[0] == 255)  
        send(sock, buf, len);  
    printf("%s", buf);  
}
```

“No send after read”

$$P = \left\{ \begin{array}{l} \{[\text{read}], [\text{read}, \text{read}], \dots\} \\ \{[\text{send}], [\text{send}, \text{read}], [\text{send}, \text{send}, \text{read}], \dots\} \\ \dots \end{array} \right\}$$

Example

Suppose that we want a simple policy

```
while(read(&buf, &len, fp)) {  
    if(buf[0] == 255)  
        send(sock, buf, len);  
    printf("%s", buf);  
}
```

“No send after read”

$$P = \left\{ \begin{array}{l} \{[read], [read, read], \dots\} \\ \{[send], [send, read], [send, send, read], \dots\} \\ \dots \end{array} \right\}$$

Does the program satisfy this policy?

Formalizing execution monitoring

Recall the key feature of EM:

- ▶ Must work by monitoring the execution of a single execution

Formalizing execution monitoring

Recall the key feature of EM:

- ▶ Must work by monitoring the execution of a single execution

We should be able to express P using simpler means

- ▶ Let $\hat{P} \subseteq \Psi$ be a set of executions
- ▶ We can define \hat{P} so that:

$$\Sigma_S \in P \iff \sigma \in \hat{P} \text{ for all } \sigma \in \Sigma_S$$

Formalizing execution monitoring

Recall the key feature of EM:

- ▶ Must work by monitoring the execution of a single execution

We should be able to express P using simpler means

- ▶ Let $\hat{P} \subseteq \Psi$ be a set of executions
- ▶ We can define \hat{P} so that:

$$\Sigma_S \in P \iff \sigma \in \hat{P} \text{ for all } \sigma \in \Sigma_S$$

Why is this simpler?

Formalizing execution monitoring

Recall the key feature of EM:

- ▶ Must work by monitoring the execution of a single execution

We should be able to express P using simpler means

- ▶ Let $\hat{P} \subseteq \Psi$ be a set of executions
- ▶ We can define \hat{P} so that:

$$\Sigma_S \in P \iff \sigma \in \hat{P} \text{ for all } \sigma \in \Sigma_S$$

Why is this simpler?

- ▶ We don't need to specify all allowed systems to get P !
- ▶ We can just specify all allowed *executions* instead.

Formalizing execution monitoring

Recall the key feature of EM:

- ▶ Must work by monitoring the execution of a single execution

We should be able to express P using simpler means

- ▶ Let $\hat{P} \subseteq \Psi$ be a set of executions
- ▶ We can define \hat{P} so that:

$$\Sigma_S \in P \iff \sigma \in \hat{P} \text{ for all } \sigma \in \Sigma_S$$

Why is this simpler?

- ▶ We don't need to specify all allowed systems to get P !
- ▶ We can just specify all allowed *executions* instead.
- ▶ Again, why is this simpler?

Example

Let's revisit the policy from before

```
while(read(&buf, &len, fp)) {  
    if(buf[0] == 255)  
        send(sock, buf, len);  
    printf("%s", buf);  
}
```

“No send after read”

Example

Let's revisit the policy from before

```
while(read(&buf, &len, fp)) {  
    if(buf[0] == 255)  
        send(sock, buf, len);  
    printf("%s", buf);  
}
```

“No send after read”

$$\hat{P} = \left\{ \begin{array}{l} [\text{read}, \text{read}, \dots], \\ [\text{send}, \text{read}, \dots], \\ [\text{send}, \text{send}, \text{read}], \\ \dots \end{array} \right\}$$

Which policies are enforceable?

We know they have to be expressible in terms of \hat{P}

Which policies are enforceable?

We know they have to be expressible in terms of \hat{P}

- ▶ i.e., policies determined by each element in Σ_S alone
- ▶ These are called **properties**

Are all properties enforceable?

Which policies are enforceable?

We know they have to be expressible in terms of \hat{P}

- ▶ i.e., policies determined by each element in Σ_S alone
- ▶ These are called **properties**

Are all properties enforceable?

- ▶ Recall: can't use information about *future* steps

Which policies are enforceable?

We know they have to be expressible in terms of \hat{P}

- ▶ i.e., policies determined by each element in Σ_S alone
- ▶ These are called **properties**

Are all properties enforceable?

- ▶ Recall: can't use information about *future* steps
- ▶ So if σ' is a **prefix** of σ and $\sigma' \notin \hat{P}$, but $\sigma \in \hat{P}$

Which policies are enforceable?

We know they have to be expressible in terms of \hat{P}

- ▶ i.e., policies determined by each element in Σ_S alone
- ▶ These are called **properties**

Are all properties enforceable?

- ▶ Recall: can't use information about *future* steps
- ▶ So if σ' is a **prefix** of σ and $\sigma' \notin \hat{P}$, but $\sigma \in \hat{P}$
- ▶ ...then P isn't enforceable

Which policies are enforceable?

We know they have to be expressible in terms of \hat{P}

- ▶ i.e., policies determined by each element in Σ_S alone
- ▶ These are called **properties**

Are all properties enforceable?

- ▶ Recall: can't use information about *future* steps
- ▶ So if σ' is a **prefix** of σ and $\sigma' \notin \hat{P}$, but $\sigma \in \hat{P}$
- ▶ ...then P isn't enforceable

Prefix closure

Enforceable policies are **prefix-closed**:

- ▶ If a trace is in \hat{P} then so are all its prefixes
- ▶ If a trace isn't in \hat{P} , then none of its extensions are

Our goal: define policies that can be enforced *for real*

Our goal: define policies that can be enforced *for real*

- ▶ Ultimately, we'd like to know if a policy violation is underway
- ▶ We have finite time to wait around for this to happen

Our goal: define policies that can be enforced *for real*

- ▶ Ultimately, we'd like to know if a policy violation is underway
- ▶ We have finite time to wait around for this to happen

Finite refutability

A property P is **finitely refutable** if whenever a trace σ is *not* in \hat{P} , there exists some *finite prefix* σ' of σ that is also not in \hat{P} .

$$\sigma \notin \hat{P} \iff \exists i. \sigma[..i] \notin \hat{P}$$

where $\sigma[..i]$ corresponds to the subsequence of σ from its beginning to position i .

Prefix closure and **finite refutability** simplify matters further:

Prefix closure and **finite refutability** simplify matters further:

- ▶ We don't need to specify *all allowed executions*

Prefix closure and **finite refutability** simplify matters further:

- ▶ We don't need to specify *all allowed executions*
- ▶ Instead, specify a set of finite prefixes

These prefixes are “bad things” disallowed by the policy

Prefix closure and **finite refutability** simplify matters further:

- ▶ We don't need to specify *all allowed executions*
- ▶ Instead, specify a set of finite prefixes

These prefixes are “bad things” disallowed by the policy

- ▶ Because our policies are properties, we look for bad things in “real time” on a single execution

Prefix closure and **finite refutability** simplify matters further:

- ▶ We don't need to specify *all allowed executions*
- ▶ Instead, specify a set of finite prefixes

These prefixes are “bad things” disallowed by the policy

- ▶ Because our policies are properties, we look for bad things in “real time” on a single execution
- ▶ By prefix closure, once we see the bad thing happen, we know the policy is permanently violated

Prefix closure and **finite refutability** simplify matters further:

- ▶ We don't need to specify *all allowed executions*
- ▶ Instead, specify a set of finite prefixes

These prefixes are “bad things” disallowed by the policy

- ▶ Because our policies are properties, we look for bad things in “real time” on a single execution
- ▶ By prefix closure, once we see the bad thing happen, we know the policy is permanently violated
- ▶ By finite refutability, if a policy violation happens we will (in principle) detect it

Enforceable policies

Properties satisfying prefix closure and finite refutability are called **safety properties**

What policies are safety properties?

Enforceable policies

Properties satisfying prefix closure and finite refutability are called **safety properties**

What policies are safety properties?

- ▶ **Access control**, defined broadly as policies that proscribe unacceptable operations. This includes filesystem permissions, bounds checking, read-xor-execute, ...

Enforceable policies

Properties satisfying prefix closure and finite refutability are called **safety properties**

What policies are safety properties?

- ▶ **Access control**, defined broadly as policies that proscribe unacceptable operations. This includes filesystem permissions, bounds checking, read-xor-execute, ...
- ▶ **Information flow** is *not* safety: it cannot be defined in terms of individual executions.

Enforceable policies

Properties satisfying prefix closure and finite refutability are called **safety properties**

What policies are safety properties?

- ▶ **Access control**, defined broadly as policies that proscribe unacceptable operations. This includes filesystem permissions, bounds checking, read-xor-execute, ...
- ▶ **Information flow** is *not* safety: it cannot be defined in terms of individual executions. Did we define information flow with “no send after read”?

Enforceable policies

Properties satisfying prefix closure and finite refutability are called **safety properties**

What policies are safety properties?

- ▶ **Access control**, defined broadly as policies that proscribe unacceptable operations. This includes filesystem permissions, bounds checking, read-xor-execute, ...
- ▶ **Information flow** is *not* safety: it cannot be defined in terms of individual executions. Did we define information flow with “no send after read”?
- ▶ **Availability** is *not* safety: any partial execution can be *extended* to grant access to the resource in question, so we can't define a set of finite prefixes to characterize availability.

Before, we enforced
“no send after read”

```
while(read(&buf, &len, fp)) {  
    if(buf[0] == 255)  
        send(sock, buf, len);  
    printf("%s", buf);  
}
```

Safety and information flow

Before, we enforced
“no send after read”

We wanted to prevent:

$fp \longrightarrow sock$

```
while(read(&buf, &len, fp)) {  
    if(buf[0] == 255)  
        send(sock, buf, len);  
    printf("%s", buf);  
}
```

Safety and information flow

Before, we enforced
“no send after read”

We wanted to prevent:

$fp \longrightarrow sock$

How is this *not* an
information flow policy?

```
while(read(&buf, &len, fp)) {  
    if(buf[0] == 255)  
        send(sock, buf, len);  
    printf("%s", buf);  
}
```

Safety and information flow

Before, we enforced
“no send after read”

We wanted to prevent:

$fp \longrightarrow sock$

How is this *not* an
information flow policy?

```
while(read(&buf, &len, fp)) {  
    if(buf[0] == 255)  
        send(sock, buf, len);  
    printf("%s", buf);  
}
```

```
while(read(&buf, &len, fp)) {  
    memset(buf, 0, len);  
    send(sock, , len);  
    printf("%s", buf);  
}
```

Does this flow fp to $sock$?

Safety and information flow

Before, we enforced
“no send after read”

We wanted to prevent:

$fp \longrightarrow sock$

How is this *not* an
information flow policy?

This policy *approximates*
information flow

- ▶ Prevents a flow from happening
- ▶ Also prevents other things

```
while(read(&buf, &len, fp)) {  
    if(buf[0] == 255)  
        send(sock, buf, len);  
    printf("%s", buf);  
}
```

```
while(read(&buf, &len, fp)) {  
    memset(buf, 0, len);  
    send(sock, , len);  
    printf("%s", buf);  
}
```

Does this flow fp to $sock$?

Information flow isn't EM-enforceable

Suppose x and y are bits

```
if (x)
  y = 0;
else
  y = 1;
```

With executions:

$$\left\{ \begin{array}{l} [(x \mapsto 0, y \mapsto 0), (x \mapsto 0, y \mapsto 1)] \\ [(x \mapsto 0, y \mapsto 1), (x \mapsto 0, y \mapsto 1)] \\ [(x \mapsto 1, y \mapsto 0), (x \mapsto 1, y \mapsto 0)] \\ [(x \mapsto 1, y \mapsto 1), (x \mapsto 1, y \mapsto 0)] \end{array} \right\}$$

Information flow isn't EM-enforceable

Suppose x and y are bits

```
if (x)
  y = 0;
else
  y = 1;
```

What *is* information flow from x to y ?

With executions:

$$\left\{ \begin{array}{l} [(x \mapsto 0, y \mapsto 0), (x \mapsto 0, y \mapsto 1)] \\ [(x \mapsto 0, y \mapsto 1), (x \mapsto 0, y \mapsto 1)] \\ [(x \mapsto 1, y \mapsto 0), (x \mapsto 1, y \mapsto 0)] \\ [(x \mapsto 1, y \mapsto 1), (x \mapsto 1, y \mapsto 0)] \end{array} \right\}$$

Information flow isn't EM-enforceable

Suppose x and y are bits

```
if (x)
  y = 0;
else
  y = 1;
```

What is information flow from x to y ?

Changes to x cause changes in y

With executions:

$$\left\{ \begin{array}{l} [(x \mapsto 0, y \mapsto 0), (x \mapsto 0, y \mapsto 1)] \\ [(x \mapsto 0, y \mapsto 1), (x \mapsto 0, y \mapsto 1)] \\ [(x \mapsto 1, y \mapsto 0), (x \mapsto 1, y \mapsto 0)] \\ [(x \mapsto 1, y \mapsto 1), (x \mapsto 1, y \mapsto 0)] \end{array} \right\}$$

A thought experiment

Let S_1 :

```
if(x)
  y = 0;
else
  y = 1;
```

And S_2 :

```
x, y = 0, 1;
```

And S_3 :

```
x, y = 1, 0;
```

A thought experiment

Let S_1 :

```
if (x)
  y = 0;
else
  y = 1;
```

Experiment as follows:

1. You pick an initial value for x .

And S_2 :

```
x, y = 0, 1;
```

And S_3 :

```
x, y = 1, 0;
```

A thought experiment

Let S_1 :

```
if (x)
  y = 0;
else
  y = 1;
```

And S_2 :

```
x, y = 0, 1;
```

And S_3 :

```
x, y = 1, 0;
```

Experiment as follows:

1. You pick an initial value for x .
2. I (secretly) pick $i \in \{1, 2, 3\}$, run S_i on your input.

A thought experiment

Let S_1 :

```
if (x)
  y = 0;
else
  y = 1;
```

And S_2 :

```
x, y = 0, 1;
```

And S_3 :

```
x, y = 1, 0;
```

Experiment as follows:

1. You pick an initial value for x .
2. I (secretly) pick $i \in \{1, 2, 3\}$, run S_i on your input.
3. You see the execution σ , try to guess i .

A thought experiment

Let S_1 :

```
if (x)
  y = 0;
else
  y = 1;
```

And S_2 :

```
x, y = 0, 1;
```

And S_3 :

```
x, y = 1, 0;
```

Experiment as follows:

1. You pick an initial value for x .
2. I (secretly) pick $i \in \{1, 2, 3\}$, run S_i on your input.
3. You see the execution σ , try to guess i .

Suppose I give you:

$$[(x \mapsto 0, y \mapsto 0), (x \mapsto 0, y \mapsto 1)]$$

How to distinguish between S_1 and S_2 ?

A (modified) thought experiment

Let S_1 :

```
if (x)
  y = 0;
else
  y = 1;
```

And S_2 :

```
x, y = 0, 1;
```

And S_3 :

```
x, y = 1, 0;
```

Experiment as follows:

1. You pick **two** initial values for x .
2. I (secretly) pick $i \in \{1, 2, 3\}$, run S_i on your inputs.
3. You see the executions σ_1, σ_2 , try to guess i .

A (modified) thought experiment

Let S_1 :

```
if (x)
  y = 0;
else
  y = 1;
```

And S_2 :

```
x, y = 0, 1;
```

And S_3 :

```
x, y = 1, 0;
```

Experiment as follows:

1. You pick **two** initial values for x .
2. I (secretly) pick $i \in \{1, 2, 3\}$, run S_i on your inputs.
3. You see the executions σ_1, σ_2 , try to guess i .

Now you win every time

Hyperproperties

Information flow is a *hyperproperty*

Hyperproperties

Information flow is a *hyperproperty*

In particular, it is *2-safety*:

- ▶ Finitely refutable over *pairs* of traces

Hyperproperties

Information flow is a *hyperproperty*

In particular, it is *2-safety*:


- ▶ Finitely refutable over *pairs* of traces

Can generalize to *k-safety*

- ▶ Lots of interesting properties...
- ▶ Quantitative privacy
- ▶ Statistical availability



Security automata



- ▶ Formal model of an execution monitor
- ▶ “Language” for specifying policies
- ▶ Corresponds to \hat{P} from before

Image credit: Lujo Bauer


Security automaton

A **security automaton** is a non-deterministic finite or infinite-state automaton defined by:

- ▶ Q : a countable set of **automaton states**
- ▶ $Q_0 \subseteq Q$: a countable set of **initial states**
- ▶ A : a countable set of **input symbols**
- ▶ $\delta : (Q \times I) \mapsto 2^Q$: a **transition function**

Security automata: semantics

Notice: no accepting states




Security automata: semantics

Notice: no accepting states

Let Q' be the *current states*

To process execution $s_1 s_2 \dots$:




Security automata: semantics

Notice: no accepting states

Let Q' be the *current states*

To process execution $s_1 s_2 \dots$:

1. Read the next input symbol s_i



Security automata: semantics


Notice: no accepting states

Let Q' be the *current states*

To process execution $s_1 s_2 \dots$:

1. Read the next input symbol s_i
2. Change Q' to

$$\bigcup_{q \in Q'} \delta(q, s_i)$$



Security automata: semantics

Notice: no accepting states

Let Q' be the *current states*


To process execution $s_1 s_2 \dots$:

1. Read the next input symbol s_i

2. Change Q' to

$$\bigcup_{q \in Q'} \delta(q, s_i)$$

3. If Q' ever becomes empty, the input is rejected



Security automata: semantics

Notice: no accepting states

Let Q' be the *current states*


To process execution $s_1 s_2 \dots$:

1. Read the next input symbol s_i

2. Change Q' to

$$\bigcup_{q \in Q'} \delta(q, s_i)$$

3. If Q' ever becomes empty, the input is rejected



An action is allowed if a transition exists for it


Security automata: semantics

Notice: no accepting states

Let Q' be the *current states*

To process execution $s_1 s_2 \dots$:

1. Read the next input symbol s_i
2. Change Q' to
$$\bigcup_{q \in Q'} \delta(q, s_i)$$
3. If Q' ever becomes empty, the input is rejected




An action is allowed if a transition exists for it

Can process both finite and infinite sequences!

Security automata: input symbols

We label edges with *transition predicates*

- ▶ Boolean-valued and total
- ▶ Domain: A




Security automata: input symbols

We label edges with *transition predicates*

- ▶ Boolean-valued and total
- ▶ Domain: A

Let p_{ij} label edge between nodes i, j

- ▶ p_{ij} specifies a **subset of** A
- ▶ $p_{ij}(s)$ is **satisfied** if s is in that subset
- ▶ e.g., $\neg\text{read}$ is satisfied by any symbol except read



Using security automata for enforcement

Security automata can be implemented to form the basis of an execution monitor

1. Initialize automaton on program/system startup
2. Before the target executes a step, generate the corresponding symbol
3. If the automaton can make a transition, let the target execute the step
4. If the automaton can't transition, terminate the target

Assumption: bounded memory

We allowed automata to have (countably) infinite states

Assumption: bounded memory

We allowed automata to have (countably) infinite states

This is necessary for recognizing certain safety properties

- ▶ Whether a prefix should be rejected might depend on every symbol in the prefix
- ▶ The amount of memory needed to remember the past grows without bound

Assumption: bounded memory

We allowed automata to have (countably) infinite states

This is necessary for recognizing certain safety properties

- ▶ Whether a prefix should be rejected might depend on every symbol in the prefix
- ▶ The amount of memory needed to remember the past grows without bound

In practice, most security policies don't need this

- ▶ Restricting the automaton to a finite set of states is probably fine for most purposes

Assumption: target control

Need ability to terminate the target on policy violation

Assumption: target control

Need ability to terminate the target on policy violation

This makes certain safety properties non-enforceable

Assumption: target control

Need ability to terminate the target on policy violation

This makes certain safety properties non-enforceable

Real-time availability

One principal cannot be denied use of a resource for more than M seconds.

Safety characterization: “Bad thing” is an unavailable interval spanning more than M seconds.

Assumption: target control

Need ability to terminate the target on policy violation

This makes certain safety properties non-enforceable

Real-time availability

One principal cannot be denied use of a resource for more than M seconds.

Safety characterization: “Bad thing” is an unavailable interval spanning more than M seconds.

Passage of time cannot be stopped!

Assumption: mechanism integrity

To correctly enforce a policy, we must assume:

- ▶ Input symbols correspond to the actual execution
- ▶ Transitions correspond to the automaton's true transition function

Assumption: mechanism integrity

To correctly enforce a policy, we must assume:

- ▶ Input symbols correspond to the actual execution
- ▶ Transitions correspond to the automaton's true transition function

If target corrupts mechanism, it can violate these assumptions

Address this with two strategies

Assumption: mechanism integrity

To correctly enforce a policy, we must assume:

- ▶ Input symbols correspond to the actual execution
- ▶ Transitions correspond to the automaton's true transition function

If target corrupts mechanism, it can violate these assumptions

Address this with two strategies

- ▶ **Isolation:** target must be unable to write to the internal representation of the automaton

Assumption: mechanism integrity

To correctly enforce a policy, we must assume:

- ▶ Input symbols correspond to the actual execution
- ▶ Transitions correspond to the automaton's true transition function

If target corrupts mechanism, it can violate these assumptions

Address this with two strategies

- ▶ **Isolation:** target must be unable to write to the internal representation of the automaton
- ▶ **Complete mediation:** make sure that all aspects of execution that might generate input symbols are covered by implementation

Enforceable Security Policies

FRED B. SCHNEIDER
Cornell University



Recognizing safety and liveness *

Bowen Alpern¹ and Fred B. Schneider²

¹ IBM T.J. Watson Research Center, P.O. Box 218, Yorktown Heights, NY 10598, USA

² Department of Computer Science, Cornell University, Ithaca, NY 14853, USA

