

# Software Foundations of Security & Privacy

## 15316 Spring 2017

### Lecture 9:

### Authentication and Authorization

Matt Fredrikson, Jean Yang  
mfredrik@cs.cmu.edu

February 14, 2017

# Access control in a nutshell




Image Credit: Butler Lampson, Martín Abadi, Michael Burrows

The reference monitor needs to answer two questions:

- ▶ *“Who said this?”* and *“Who is trusted to access this?”*
- ▶ **Authentication:** obtaining the source of the request
- ▶ **Authorization:** interpreting and deciding the access rule

Requests arrive over a **channel**

- ▶ Network connection
- ▶ Pipe/IPC
- ▶ System call
- ▶ User input

Monitor must associate a **principal** to the channel

- ▶ **Centralized:** Requester, monitor, and policy reside on the same system under control of a single authority
- ▶ **Distributed:** Requester, monitor, and policy reside on multiple systems, with multiple authorities at differing levels of trust

**Authentication:** how do we name and identify principals?

- ▶ Login names, user IDs, process IDs, ...
- ▶ System maps process and login contexts to principals

**Authorization:** how do we decide who gets access?

- ▶ Access control list
- ▶ Stored in reference monitor, or with the object
- ▶ Maps principals to access rights

For the most part, this is a solved problem

- ▶ Just need to trust the operating system
- ▶ System implements all channels, maintains the policy
- ▶ Knows who is responsible for each request

Things are harder in a distributed system

- ▶ **Autonomy:** Path between the object and requesting principal may involve machines with differing levels of trust. We might want policies that take this into account.
- ▶ **Size & heterogeneity:** Number of users, objects, and authorities may be large, and have inconsistent interfaces and security models. Must scale to meet the demand.
- ▶ **Fault tolerance:** Remote parts of the system may become inaccessible, but we'd like to maintain as much functionality as possible.

Mathematical formalism for authentication and authorization

All systems make assumptions about authority and trust

- ▶ Logic forces us to state these assumptions precisely
- ▶ Provides sound rules for working out the consequences

To do this, we'll make some assumptions:

- ▶ Correctness of underlying operating system
- ▶ Hardware implementation is correct and secure
- ▶ Secure encryption primitives

Logic due to Reiter, influenced by Lamport, Abadi, Burrows

## **Principals** exist to make **statements**

mfredrik says “syecom *speaks for* 15316-spring17-staff”

- ▶ “syecom *speaks for* 15316-spring17-staff” is a statement
- ▶ mfredrik, syecom, 15316-spring17-staff are principals

Anything that a principal says is a statement

- ▶ requests, delegations, trust relationships, ...

Objects are the entities protected by authorization requests

# Principals

Principals are named entities, groups, and channels

- ▶ People (mfredrik, jyang2, syeom)
- ▶ Machines (linux2.cs.cmu.edu, ...)
- ▶ Groups (coursestaff, students, administrators)
- ▶ Channels (128.2.220.63, AES key #574897)

Formally, a principal is either an identifier or a key:

$$p ::= \mathbf{key}(s) \mid \mathit{identifier} \mid p.s$$

where  $s$  is a string



# Statements

Statements, where  $s$  is a string and  $p$  is a principal

$$\begin{array}{lcl} \phi & ::= & \mathbf{action}(s) \\ & | & p \mathbf{says} \phi \\ & | & p \mathbf{speaksfor} p \\ & | & s \mathbf{signed} \phi \\ & | & \mathbf{delegates}(p, p, s) \\ & | & \phi \rightarrow \phi \\ & | & \phi \wedge \phi \end{array}$$

What does this correspond to in reality?

A **digital signature scheme** is a triple  $\langle G, S, V \rangle$

- ▶ The **key generator**  $G$  takes a key length  $n$  and outputs a public/private key pair  $(pk, sk)$
- ▶ The **signing algorithm**  $S$  takes a private key  $sk$  and message  $m$ , and outputs a signature  $\sigma$ :

$$\sigma \leftarrow S(sk, m)$$

- ▶ The **verifier**  $V$  takes a public key  $pk$ , message  $m$ , and signature  $\sigma$ , and outputs either 0 or 1:

$$V(pk, m, \sigma) \mapsto \{0, 1\}$$

# Signatures: correctness

## Correctness of a signature scheme

$\langle G, S, V \rangle$  is a correct digital signature scheme if with all but negligible probability over the key pairs  $(ps, sk)$  output by  $G$ , it holds that:

$$V(pk, m, S(sk, m)) = 1$$

In other words, the verifier accepts valid signatures.

Is this enough?

# Signatures: security

Let  $\Pi = \langle G, S, V \rangle$  be a signature scheme

Define the following experiment  $\text{Forge}(\Pi, \mathcal{A}, n)$ :

1. Run  $(pk, sk) \leftarrow G(n)$ .
2. Give  $\mathcal{A}$   $pk$ , and let it run  $S(sk, \cdot)$ . Let  $Q$  be the set of queries  $\mathcal{A}$  gives to  $S(sk, \cdot)$ .
3.  $\mathcal{A}$  then outputs  $(m, \sigma)$ .
4.  $\mathcal{A}$  wins when both  $V(pk, m, \sigma) = 1$  and  $m \notin Q$ .

## Security of a signature scheme

$\Pi$  is secure if for all polynomial-time adversaries  $\mathcal{A}$  and key lengths  $n$ ,

$$\Pr[\text{Forge}(\Pi, \mathcal{A}, n) \text{ wins}] \leq \text{negl}(n)$$

for some negligible function  $\text{negl}$ .

# Signatures: summary

Signature schemes are widely used in authorization

- ▶ Pay attention to the way that security is carefully defined
- ▶ Gives a rigorous justification for  $s$  **signed**  $\phi$

This isn't a class about cryptography

- ▶ We won't cover signature schemes in more detail
- ▶ You aren't expected to memorize this definition
- ▶ The details of the scheme aren't important to us
- ▶ As long as we have one that satisfies security

# Associating trust: **speaksfor**

*A* **speaksfor** *B*

When *A* says something, we believe that *B* says it too

- ▶ 128.2.220.63 **speaksfor** mfredrik
- ▶ AES key #574897 **speaksfor** jyang2
- ▶ syeom **speaksfor** coursestaff

**speaksfor** formalizes indirection for statements

- ▶ Some principals can't communicate directly with others
- ▶ Principals often have several others speaking for them
- ▶ Roles may have a rotating set of principals speak for them

# Provable statements

Write **judgements** to denote statements that are provably true

$$\vdash s$$

This means: **statement  $s$  is provable without assumptions**

$$P \vdash s$$

This means:  **$s$  is provable using assumptions given in  $P$**

We use **inference rules** to prove things about statements

## says Introduction 1 (Says-I1)

$$\frac{s \text{ signed } \phi}{\mathbf{key}(s) \text{ says } \phi}$$

Intuitively, this rule:

- ▶ Creates a principal **key**(*s*) from a *key string s*
- ▶ Establishes **key**(*s*) said  $\phi$  given the appropriate signature



## says Introduction 2 (Says-I2)

$$\frac{\phi}{p \textbf{says } \phi}$$

What does this rule say?

- ▶ If  $\phi$  is established to be true, then  $p$  says it
- ▶ Maybe  $p$  didn't say it, but we can proceed as though it did
- ▶ Basically, principals will say true things

## **says** Introduction 3 (Says-I3)

$$\frac{p \textbf{ says } (p.s \textbf{ says } \phi)}{p.s \textbf{ says } \phi}$$

This rule might seem counterintuitive

- ▶  $p.s$  is the “principal that  $p$  calls  $s$ ”.
- ▶  $p$  can name  $s$  to be whomever it wants
- ▶  $p$  can just find someone who says  $\phi$ , and call that person  $s$
- ▶ We must accept what  $p$  says about the person it calls  $s$

## **says** Implication (Says-Impl)

$$\frac{p \textbf{ says } (\phi_1 \rightarrow \phi_2) \quad p \textbf{ says } \phi_1}{p \textbf{ says } \phi_2}$$

Recall modus ponens from propositional calculus

- ▶ Says-Impl is modus ponens over **says**
- ▶ We take principals at their word
- ▶ To the logical conclusion

## **speaksfor** Elimination 1 (Speaksfor-E1)

$$\frac{p_1 \text{ **says** } (p_2 \text{ **speaksfor** } p_1) \quad p_2 \text{ **says** } \phi}{p_1 \text{ **says** } \phi}$$

## **speaksfor** Elimination 2 (Speaksfor-E2)

$$\frac{p_1 \text{ **says** } (p_2 \text{ **speaksfor** } p_1.s) \quad p_2 \text{ **says** } \phi}{p_1.s \text{ **says** } \phi}$$

These rules deal with broad delegations of authority

- ▶ When  $p_1$  says  $p_2$  speaks for her...
- ▶ Then anything  $p_2$  says is attributable to  $p_1$

## **delegates** Elimination (Delegate-E)

$$\frac{p_1 \text{ says delegates}(p_1, p_2, s) \quad p_2 \text{ says action}(s)}{p_1 \text{ says action}(s)}$$

This rule allows more fine-grained delegation of authority

- ▶ **delegates** allows  $p_1$  to let  $p_2$  speak-for her
- ▶ But only with regard to selected actions

## Example: Certification authorities

Let's use the rules to reason about a certification authority

A certification authority is a named principal  $CA$

For our purposes, it issues statements of the form:

$K_{CA}$  **signed** (**key**( $K_A$ ) **speaksfor** **key**( $K_{CA}$ ). $A$ )

This statement is called a **certificate**

- ▶ Usually,  $K_{CA}$  is a public key known to everyone
- ▶ It could also be a symmetric key
- ▶ If so, need to ensure that  $K_{CA}$  not used as public identifier
- ▶ In practice, use secure hash functions to do this

# Example: Proof

Suppose we have:

1.  $K_{CA}$  **signed** (**key**( $K_A$ ) **speaksfor** **key**( $K_{CA}$ ). $A$ )
2.  $K_A$  **signed** **action**(**read**, **foo.txt**)

We want to derive:

**key**( $K_{CA}$ ). $A$  **says** **action**(**read**, **foo.txt**)

Proof:

3. **key**( $K_{CA}$ ) **says** (**key**( $K_A$ ) **speaksfor** **key**( $K_{CA}$ ). $A$ )  
(Says-I3 on 1)
4. **key**( $K_A$ ) **says** **action**(**read**, **foo.txt**) (Says-I3 on 2)
5. **key**( $K_{CA}$ ). $A$  **says** **action**(**read**, **foo.txt**) (Speaksfor-E2)

Example due to Mike Reiter

# Authenticating requests

Suppose the reference monitor receives a request:

$p$  **says action**( $s$ )

The monitor has a policy that enumerates:

1. A set of principals  $P$
2. The subset of principals  $P_s \subseteq P$  authorized to perform  $s$

The reference monitor needs a proof that:

$p_a$  **says action**( $s$ ), for some  $p_a \in P_s$



# Authenticating requests

There are three basic approaches for doing this.

1. **Push:** The sender of the request collects certificates necessary to prove  $p_a$  **says action**( $s$ ), and sends them with the request. The monitor then finds a proof, if one exists.
2. **Pull:** The reference monitor searches for a set of certificates sufficient to prove  $p_a$  **says action**( $s$ ), and constructs the proof.
3. **Proof-carrying:** The sender of the request collects the necessary certificates and constructs the proof of  $p_a$  **says action**( $s$ ) itself. This might be more efficient than having the server construct every proof.

## Example

$$\frac{\text{mfredrik says} \quad \frac{\text{sam speaksfor staff} \quad \text{sam says action}(s)}{\text{staff says action}(s)}}{\text{delegates}(\text{mfredrik}, \text{staff}, s)} \quad \text{mfredrik says action}(s)$$

# Centralized access control list

$$\frac{\text{mfredrik says} \quad \frac{(\text{sam speaksfor staff}) \quad \text{sam says action}(s)}{\text{staff says action}(s)}}{\text{delegates}(\text{mfredrik}, \text{staff}, s)} \quad \text{mfredrik says action}(s)$$

In a traditional access control list implementation

- ▶ The highlighted parts are stored in the reference monitor
- ▶ They're part of the TCB, and not cryptographically signed

# Pull authentication

$$\frac{\text{mfredrik says } ( \text{sam speaksfor staff } ) \quad \text{sam says action}(s)}{\text{mfredrik says } \text{delegates}(\text{mfredrik}, \text{staff}, s) \quad \text{staff says action}(s)} \\ \text{mfredrik says action}(s)$$

In pull-authentication scheme

- ▶ The red-highlighted parts are retrieved by the monitor
- ▶ The blue-highlighted part is sent by the requester
- ▶ The rest is computed by the monitor

# Push authentication

$$\frac{\text{mfredrik says} \quad \frac{\text{delegates}(\text{mfredrik}, \text{staff}, s) \quad \text{staff says action}(s)}{\text{mfredrik says action}(s)}}{(\text{sam speaksfor staff}) \quad \text{sam says action}(s)}$$

In push-authentication scheme

- ▶ The red-highlighted parts are sent by the requester
- ▶ The rest is computed by the monitor

# Proof-carrying authentication

$$\frac{\text{mfredrik says } ( \text{sam speaksfor staff } ) \quad \text{sam says action}(s)}{\text{delegates}(\text{mfredrik}, \text{staff}, s) \quad \text{staff says action}(s)} \quad \text{mfredrik says action}(s)$$

In push-authentication scheme

- ▶ The red-highlighted parts are sent by the requester
- ▶ The conclusion is **verified** by the monitor
- ▶ If the proof checks out, authorization is granted

# Next lecture

- ▶ More authorization logic
- ▶ Application to secure boot, TLS authentication
- ▶ Revocation
- ▶ Extensions to the logic