

Software Foundations of Security & Privacy

15315 Spring 2017

Lecture 10:

Authentication and Authorization

Matt Fredrikson, Jean Yang
mfredrik@cs.cmu.edu

February 16, 2017

Review: Authorization Logic

Formally, a principal is either an identifier or a key:

$$p ::= \mathbf{key}(s) \mid \mathit{identifier} \mid p.s$$

where s is a string

Statements, where s is a string and p is a principal

$$\begin{aligned} \phi ::= & \mathbf{action}(s) \\ & \mid p \mathbf{says} \phi \\ & \mid p \mathbf{speaksfor} p \\ & \mid s \mathbf{signed} \phi \\ & \mid \mathbf{delegates}(p, p, s) \\ & \mid \phi \rightarrow \phi \\ & \mid \phi \wedge \phi \end{aligned}$$

Review: inference rules

says Introduction 1 (Says-I1)

$$\frac{s \text{ signed } \phi}{\mathbf{key}(s) \text{ says } \phi}$$

Intuitively, this rule:

- ▶ Creates a principal **key**(*s*) from a *key string* *s*
- ▶ Establishes **key**(*s*) said ϕ given the appropriate signature

Review: inference rules

says Introduction 2 (Says-I2)

$$\frac{\phi}{p \text{ **says** } \phi}$$

What does this rule say?

- ▶ If ϕ is established to be true, then p says it
- ▶ Maybe p didn't say it, but we can proceed as though it did
- ▶ Basically, principals will say true things

Review: inference rules

says Introduction 3 (Says-I3)

$$\frac{p \text{ **says** } (p.s \text{ **says** } \phi)}{p.s \text{ **says** } \phi}$$

This rule might seem counterintuitive

- ▶ $p.s$ is the “principal that p calls s ”.
- ▶ p can name s to be whomever it wants
- ▶ p can just find someone who says ϕ , and call that person s
- ▶ We must accept what p says about the person it calls s

Review: inference rules

says Implication (Says-Impl)

$$\frac{p \textbf{says } (\phi_1 \rightarrow \phi_2) \quad p \textbf{says } \phi_1}{p \textbf{says } \phi_2}$$

Recall modus ponens from propositional calculus

- ▶ Says-Impl is modus ponens over **says**
- ▶ We take principals at their word
- ▶ To the logical conclusion

Review: inference rules

speaksfor Elimination 1 (Speaksfor-E1)

$$\frac{p_1 \text{ **says** } (p_2 \text{ **speaksfor** } p_1) \quad p_2 \text{ **says** } \phi}{p_1 \text{ **says** } \phi}$$

speaksfor Elimination 2 (Speaksfor-E2)

$$\frac{p_1 \text{ **says** } (p_2 \text{ **speaksfor** } p_1.s) \quad p_2 \text{ **says** } \phi}{p_1.s \text{ **says** } \phi}$$

These rules deal with broad delegations of authority

- When p_1 says p_2 speaks for her...

Then anything p_2 says is attributable to p_1

Review: inference rules

delegates Elimination (Delegate-E)

$$\frac{p_1 \text{ says delegates}(p_1, p_2, s) \quad p_2 \text{ says action}(s)}{p_1 \text{ says action}(s)}$$

This rule allows more fine-grained delegation of authority

- ▶ **delegates** allows p_1 to let p_2 speak-for her
- ▶ But only with regard to selected actions

Roles

In some cases, principals want to limit their authority

- ▶ Running untrusted code
- ▶ Performing mundane, non-critical tasks

Want to model principals acting **as** other principals

In this logic, we can issue statements of the form:

$$p.role \textbf{says } \phi$$

A role is just another principal named by us

Example: programs as roles

If principal p wants to run program a with text t

- ▶ Could emit $p.t$ **says** ϕ when program requests ϕ

Alternatively, compute a secure hash h of t

- ▶ Then state p **says** (h **speaksfor** t)
- ▶ What's the benefit of doing it this way?

Reference monitors don't typically want to track code hashes

- ▶ Instead, refer to named apps (e.g., "Chrome")
- ▶ State that hash values speak for the named app
- ▶ Have CA or trusted developer sign this statement

Example due to Mike Reiter

Example: authorizing a program

Suppose we want to execute program a on channel c

1. Read the code t of program a from disk
2. Compute a secure hash h of t
3. Reference monitor attempts to prove that h **speaksfor** a
4. Create process d , copy t into d
5. Emit p **says** c **speaksfor** $p.a$
6. Give d access to channel c

d can now send requests on c

If $p.a$ is authorized for a request, it's granted

Example due to Mike Reiter

Example: malicious code

Malicious code often alters program code on disk

In our example, this means that:

- ▶ Altered code t' creates new hash h'
- ▶ With high probability, $h \neq h'$!
- ▶ Loader can't prove h' **speaksfor** a

Can utilize CA to say which programs are safe

- ▶ K_{CA} **signed** (Chrome **speaksfor key**(K_{CA}).trustedSW)
- ▶ trustedSW is a group name
- ▶ Authorization rights can be established for group
- ▶ Based on trust in CA

Example due to Mike Reiter

Secure boot

Idea: establish a group `trustedhosts`

- ▶ OS is validated before booting
- ▶ OS validates all programs before running them

Validating OS is much the same as validating any other program

- ▶ $h(OS)$ **speaksfor** OS
- ▶ $h(OS)$ is stored in the system, checked against code before boot

Host (N) can set up channel for OS to use:

1. Generate signing keys $pk_{N.OS}, sk_{N.OS}$
2. Give OS: $sk_{N.OS}$ and pk_N **signed key**($pk_{N.OS}$) **speaksfor key**(pk_N).OS

Slide from Mike Reiter

Trusted platform

Effort developed by Trusted Computing Group (TCG)

- ▶ Industry consortium with about 100 members founded in 2003
- ▶ Standardize a *trusted platform*, using principles we've discussed
- ▶ Technical specs relate to hardware, BIOS, and OS features

Some of the main goals of a TP:

- ▶ **Software attestation:** users can validate software running on platform
- ▶ **Binding:** mechanism for secure storage of crypto keys and encryption
- ▶ **Sealing:** specify conditions under which decryption can occur

Slide from Mike Reiter

Trusted platform module (TPM)

Basis for trusted platform is a hardware component called the TPM

- ▶ Restricted software interface, isolated from attack “from above”
- ▶ (Somewhat) resistant to hardware tampering
- ▶ Initialized with a unique key pair pk_{TPM}, sk_{TPM}

The manufacturer loads a certificate onto each TPM:

pk_{TPME} **signed** (**key**(pk_{TPM}) **speaksfor** **key**(pk_{TPME}).*TrustedModules*)

TrustedModules is a group

- ▶ All TPMs manufactured by that entity
- ▶ Trust in TPM-signed statements reduces to trust in the manufacturer

Slide from Mike Reiter

Roots of trust

Background: in TPM-land, you'll often hear of “measurements”

- ▶ A “measurement” of a piece of software is a hash of its code
- ▶ Why not just say “code hash”?

Root of trust for measurement

- ▶ Component that begins measurement of running software
- ▶ Usually, the first thing to run on boot
- ▶ Can run later, but everything before it must be trusted

Root of trust for reporting

- ▶ Component responsible for storing measurements as they are taken
- ▶ Must also prevent tampering or undoing of stored measurements

Platform configuration registers


TPM contains 16 20-byte PCR

PCRs hold measurements of running software

- ▶ Updated before software is run
- ▶ Can't be reset until reboot

PCRs *extended* to reflect newly-loaded code

Provide a record of all running software



Slide from Mike Reiter

Authenticated boot

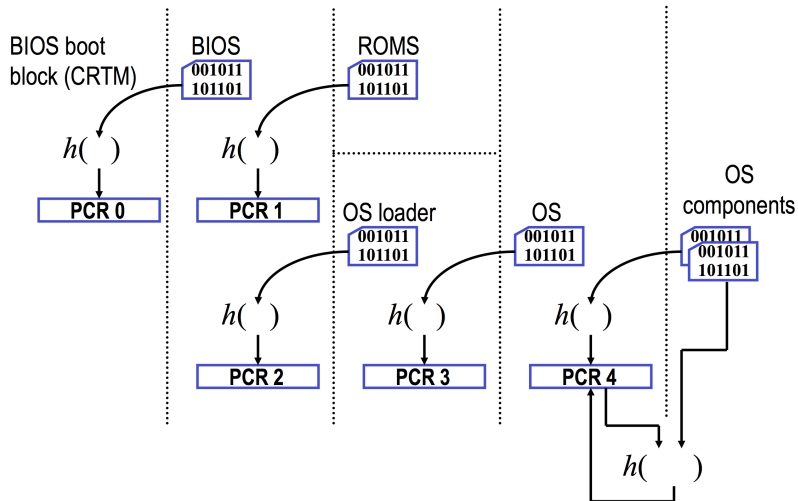


Image credit: Mike Reiter

Remote attestation

Goal: convince a remote machine that we're running the right code

TPM can provide a signed copy of the PCR values


pk_{TPM} **signed** pcr_vals

TPM will also respond with records of what the measurements summarize

- ▶ Name, version number, etc. for loaded software
- ▶ Order in which it was loaded
- ▶ Enough for remote machine to check PCR values

Any issues with this setup?

Example



1. pk_{CA} **signed** (**key**(pk_N) **speaksfor** **key**(pk_{CA}). N)
2. pk_{CA} **signed** (**key**(pk_U) **speaksfor** **key**(pk_{CA}). U)
3. pk_N **signed** (**key**($pk_{N.OS}$) **speaksfor** **key**(pk_{CA}). $N.OS$)
4. $pk_{N.OS}$ **signed** (**key**($pk_{N.OS}$). U **says** F)


Can we prove **key**(pk_{CA}). U **says** F ?

Example: proof

1. **key**(pk_{CA}) **says** (**key**(pk_N) **speaksfor** **key**(pk_{CA}). N) (Says-I1)
2. **key**(pk_{CA}) **says** (**key**(pk_U) **speaksfor** **key**(pk_{CA}). U) (Says-I1)
3. **key**(pk_N) **says** (**key**($pk_{N.OS}$) **speaksfor** **key**(pk_{CA}). $N.OS$) (Says-I1)
4. **key**($pk_{N.OS}$) **says** (**key**($pk_{N.OS}$). U **says** F) (Says-I1)
5. **key**($pk_{N.OS}$). U **says** F (Says-I3)
6. **key**(pk_{CA}). N **says** (**key**($pk_{N.OS}$) **speaksfor** **key**(pk_{CA}). $N.OS$) (Speaksfor-E2)
7. **key**(pk_{CA}). $N.OS$ **says** (**key**($pk_{N.OS}$). U **says** F) (Speaksfor-E2)

Can't make any more progress!


Example



1. pk_{CA} **signed** (**key**(pk_N) **speaksfor** **key**(pk_{CA}). N)
2. pk_{CA} **signed** (**key**(pk_U) **speaksfor** **key**(pk_{CA}). U)
3. pk_N **signed** (**key**($pk_{N.OS}$) **speaksfor** **key**(pk_{CA}). $N.OS$)
4. $pk_{N.OS}$ **signed** (**key**($pk_{N.OS}$). U **says** F)

What went wrong?

Example



1. pk_{CA} **signed** (**key**(pk_N) **speaksfor** **key**(pk_{CA}). N)
2. pk_{CA} **signed** (**key**(pk_U) **speaksfor** **key**(pk_{CA}). U)
3. pk_N **signed** (**key**($pk_{N.OS}$) **speaksfor** **key**(pk_{CA}). $N.OS$)
4. $pk_{N.OS}$ **signed** (**key**($pk_{N.OS}$). U **says** F)
5. pk_U **signed** (**key**(pk_{CA}). $N.OS.U$ **speaksfor** **key**(pk_{CA}). U)


User needs to delegate authority to “who the OS calls U”

Example: proof

1. **key(pk_{CA}) says (key(pk_N) speaksfor key(pk_{CA}). N)** (Says-I1)
2. **key(pk_{CA}) says (key(pk_U) speaksfor key(pk_{CA}). U)** (Says-I1)
3. **key(pk_N) says (key($pk_{N.OS}$) speaksfor key(pk_{CA}). $N.OS$)** (Says-I1)
4. **key($pk_{N.OS}$) says (key($pk_{N.OS}$). U says F)** (Says-I1)
5. **key($pk_{N.OS}$). U says F** (Says-I3)
6. **key(pk_{CA}). N says (key($pk_{N.OS}$) speaksfor key(pk_{CA}). $N.OS$)**
(Speaksfor-E2)
7. **key(pk_{CA}). $N.OS$ says (key($pk_{N.OS}$). U says F)** (Speaksfor-E2)
8. **key(pk_U) says (key(pk_{CA}). $N.OS.U$ speaksfor key(pk_{CA}). U)** (Says-I1)
9. **key(pk_{CA}). U says (key(pk_{CA}). $N.OS.U$ speaksfor key(pk_{CA}). U)**
(Speaksfor-E2)

Stuck again!


Example



1. pk_{CA} **signed** (**key**(pk_N) **speaksfor** **key**(pk_{CA}). N)
2. pk_{CA} **signed** (**key**(pk_U) **speaksfor** **key**(pk_{CA}). U)
3. pk_N **signed** (**key**($pk_{N.OS}$) **speaksfor** **key**(pk_{CA}). $N.OS$)
4. $pk_{N.OS}$ **signed** (**key**($pk_{N.OS}$). U **says** F)
5. pk_U **signed** (**key**(pk_{CA}). $N.OS.U$ **speaksfor** **key**(pk_{CA}). U)

Still missing something...

Example



1. pk_{CA} **signed** (**key**(pk_N) **speaksfor** **key**(pk_{CA}). N)
2. pk_{CA} **signed** (**key**(pk_U) **speaksfor** **key**(pk_{CA}). U)
3. pk_N **signed** (**key**($pk_{N.OS}$) **speaksfor** **key**(pk_{CA}). $N.OS$)
4. $pk_{N.OS}$ **signed** (**key**($pk_{N.OS}$). U **says** F)
5. pk_U **signed** (**key**(pk_{CA}). $N.OS.U$ **speaksfor** **key**(pk_{CA}). U)
6. $pk_{N.OS}$ **signed** (**key**($pk_{N.OS}$). U **speaksfor** **key**(pk_{CA}). $N.OS.U$)

OS needs to vouch for authenticity of “who the OS calls U”


Example: proof

1. **key**(pk_{CA}) **says** (**key**(pk_N) **speaksfor** **key**(pk_{CA}). N) (Says-I1)
2. **key**(pk_{CA}) **says** (**key**(pk_U) **speaksfor** **key**(pk_{CA}). U) (Says-I1)
3. **key**(pk_N) **says** (**key**($pk_{N.OS}$) **speaksfor** **key**(pk_{CA}). $N.OS$) (Says-I1)
4. **key**($pk_{N.OS}$) **says** (**key**($pk_{N.OS}$). U **says** F) (Says-I1)
5. **key**($pk_{N.OS}$). U **says** F (Says-I3)
6. **key**(pk_{CA}). N **says** (**key**($pk_{N.OS}$) **speaksfor** **key**(pk_{CA}). $N.OS$) (Speaksfor-E2)
7. **key**(pk_{CA}). $N.OS$ **says** (**key**($pk_{N.OS}$). U **says** F) (Speaksfor-E1)
8. **key**(pk_U) **says** (**key**(pk_{CA}). $N.OS.U$ **speaksfor** **key**(pk_{CA}). U) (Says-I1)
9. **key**(pk_{CA}). U **says** (**key**(pk_{CA}). $N.OS.U$ **speaksfor** **key**(pk_{CA}). U) (Speaksfor-E2)
10. **key**($pk_{N.OS}$) **says** (**key**($pk_{N.OS}$). U **speaksfor** **key**(pk_{CA}). $N.OS.U$)

Example: proof continued

5. **key**($pk_{N.OS}$). U **says** F (Says-I3)
6. **key**(pk_{CA}). N **says** (**key**($pk_{N.OS}$) **speaksfor** **key**(pk_{CA}). $N.OS$)
(Speaksfor-E2)
7. **key**(pk_{CA}). $N.OS$ **says** (**key**($pk_{N.OS}$). U **says** F) (Speaksfor-E2)
8. **key**(pk_U) **says** (**key**(pk_{CA}). $N.OS.U$ **speaksfor** **key**(pk_{CA}). U) (Says-I1)
9. **key**(pk_{CA}). U **says** (**key**(pk_{CA}). $N.OS.U$ **speaksfor** **key**(pk_{CA}). U)
(Speaksfor-E2)
10. **key**($pk_{N.OS}$) **says** (**key**($pk_{N.OS}$). U **speaksfor** **key**(pk_{CA}). $N.OS.U$)
11. **key**(pk_{CA}). $N.OS$ **says** (**key**($pk_{N.OS}$). U **speaksfor** **key**(pk_{CA}). $N.OS.U$)
(Speaksfor-E2)
12. **key**(pk_{CA}). $N.OS.U$ **says** F (Speaksfor-E2)
13. **key**(pk_{CA}). U **says** F (Speaksfor-E1)

Revocation

 Carnegie Mellon University [US] | <https://www.cmu.edu>

When you see this, it means that CMU's server sent a certificate

pk_{CA} **signed** (**key**($pk'_{cmu.edu'}$) **speaksfor** **key**(pk_{CA}).' $cmu.edu'$)

What happens if $sk'_{cmu.edu'}$ is compromised?

- ▶ CA's certificate becomes a lie!
- ▶ To maintain trust, CA may need to **revoke** the certificate

Countersigning

Idea: rely on another authority O to sign for validity of the certificate

Now the CA releases a *weaker* certificate

$$pk_{CA} \text{ signed } (\text{key}(pk_O) \text{ says key}(pk_{cmu.edu'}) \text{ speaksfor key}(pk_O.'cmu.edu'))$$
$$\rightarrow (\text{key}(pk_{cmu.edu'}) \text{ speaksfor key}(pk_{CA}).'cmu.edu'))$$

And O countersigns:

$$pk_O \text{ signed}$$
$$(\text{date}() < '2017.2.17' \rightarrow \text{key}(pk_{cmu.edu'}) \text{ speaksfor key}(pk_O).'cmu.edu'))$$

Slide from Mike Reiter

Example

Assumptions:

1. pk_{CA} **signed** ($\text{key}(pk_O)$ **says** $\text{key}(pk_A)$ **speaksfor** $\text{key}(pk_O.A) \rightarrow (\text{key}(pk_A)$ **speaksfor** $\text{key}(pk_{CA}).A)$)
 2. pk_O **signed** ($date < '2017.2.17' \rightarrow \text{key}(pk_A)$ **speaksfor** $\text{key}(pk_O).A$)
 3. $date < '2017.2.17'$
-
1. $\text{key}(pk_{CA})$ **says** ($\text{key}(pk_O)$ **says** $\text{key}(pk_A)$ **speaksfor** $\text{key}(pk_O.A) \rightarrow (\text{key}(pk_A)$ **speaksfor** $\text{key}(pk_{CA}).A)$)
 2. $\text{key}(pk_O)$ **says** ($date < '2017.2.17' \rightarrow \text{key}(pk_A)$ **speaksfor** $\text{key}(pk_O).A$)
 3. $\text{key}(pk_O)$ **says** $date < '2017.2.17'$
 4. $\text{key}(pk_O)$ **says** $\text{key}(pk_A)$ **speaksfor** $\text{key}(pk_O).A$
 5. $\text{key}(pk_{CA})$ **says** ($\text{key}(pk_O)$ **says** $\text{key}(pk_A)$ **speaksfor** $\text{key}(pk_O).A$)
 6. $\text{key}(pk_{CA})$ **says** ($\text{key}(pk_A)$ **speaksfor** $\text{key}(pk_{CA}).A$)

Why countersign?

Why doesn't `cmu.edu` just obtain a new certificate with each connection?

- ▶ Efficiency and scalability, perhaps
- ▶ CA's private signing key would need to be kept on an online server
- ▶ This exposes it to all kinds of risk, better to keep it offline

Does countersigning address the risk?

- ▶ If O 's signing key becomes compromised, what happens?
- ▶ Can extend the period where the compromised certificate is accepted
- ▶ But O cannot issue new certificates, so the damage is limited

Further reading

B. Lampson, M. Abadi, M. Burrows. *Authorization in Distributed Systems*. In ACM TOCS, 1992.

M. Abadi. *Logic in Access Control*. In Logic in Computer Science, 2003.

L. Bauer, S. Garriss, M. K. Reiter. *Efficient Proving for Practical Distributed Access-Control Systems*. In ESORICS, 2007.