

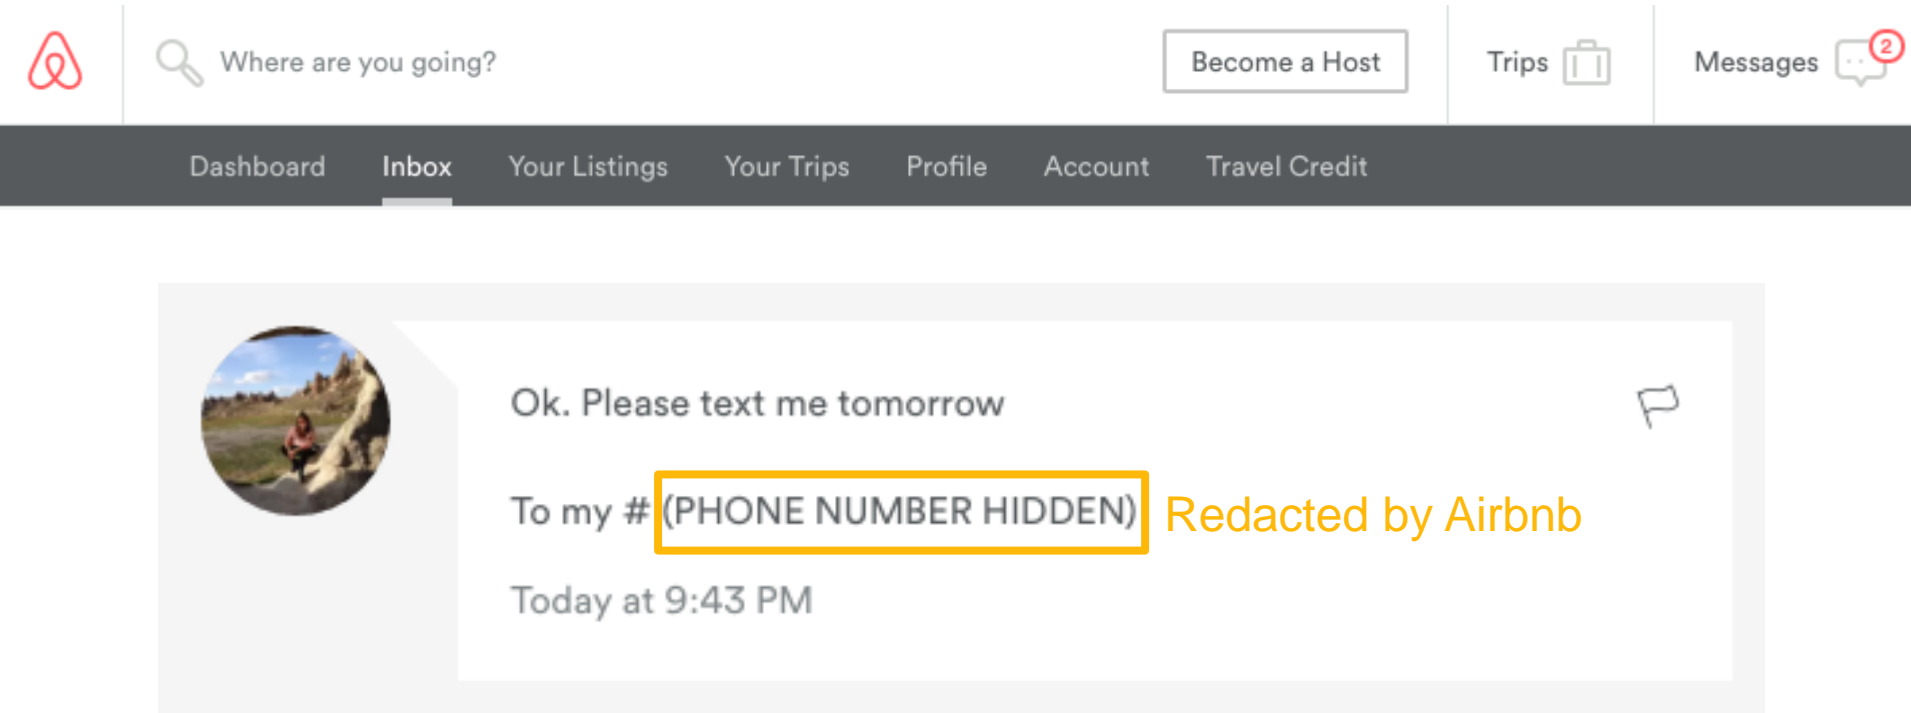
Software Foundations of Security and
Privacy (15-316, spring 2017)
Lecture 11: Information Flow (1)

Jean Yang

jyang2@andrew.cmu.edu

Goal: Keep Secrets Secret

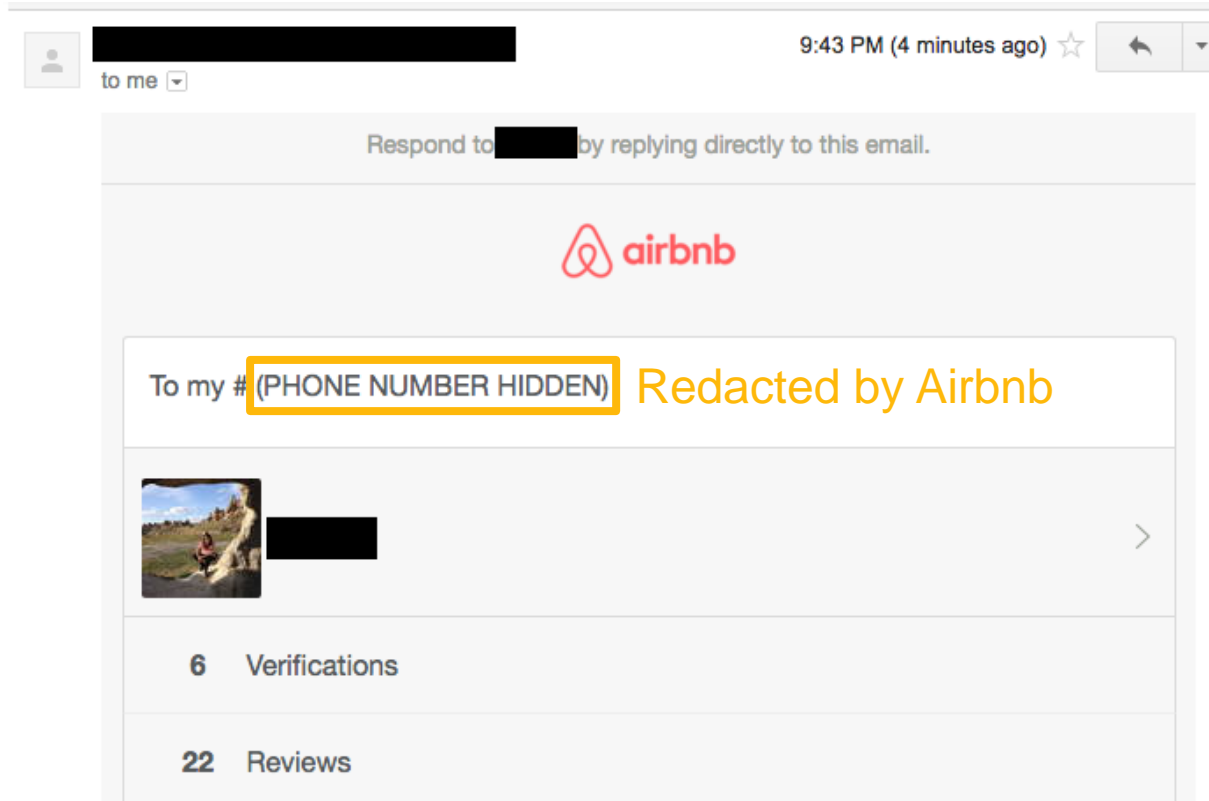
Example courtesy of Chelsea Voss



Airbnb has a policy of blocking phone numbers so communications happen through their application.

Redacting Phone Numbers...

Example courtesy of Chelsea Voss



Phone number remains redacted in email view.

Missed a spot!

Example courtesy of Chelsea Voss


The screenshot shows the Airbnb website's header and a message preview. The header includes the Airbnb logo, a search bar with the text "Where are you going?", a "Become a Host" button, and links for "Trips", "Messages" (with a red notification badge showing "2"), and "Help". Below the header is a navigation bar with links: "Dashboard", "Inbox" (highlighted), "Your Listings", "Your Trips", "Profile", "Account", and "Travel Credit".

A light blue promotional banner is visible, stating "Earn \$100 travel credit" and "Give your friends \$20 off their first trip on Airbnb and you'll get up to \$100 travel credit." It includes "Invite Friends" and "Later" buttons.

Below the banner is a message preview section. It shows a dropdown menu for "All Messages (2)". The message preview includes a circular profile picture, a redacted name, a timestamp of "9:43 PM", a redacted phone number, and a message body that says "To my # [redacted] Washington, DC (Mar 15 - 16, 2016)". To the right of the message body, it says "Accepted" in green and "\$115". A red arrow points from the text "Actual phone number! Redacted by me and not Airbnb." to the redacted phone number in the message preview.

Phone number is visible in message preview.


Main Takeaway



If companies can't even protect information when their own financial interests are at stake, then we should be really afraid.

Alternative Takeaway

Companies don't have the tools to prevent unauthorized information flows even when they are motivated to do so!



This Lecture: A Tribute to Max Krohn



@maxtaco on Twitter.

- Founded Thespark.com, OKCupid, and Keybase.
- Built OKWS for OKCupid as a PhD student at MIT.
- Continued using his research to make OKCupid's backend better throughout his PhD.



Part One: What's Wrong with Access Control?

Problems with Access Control

- Need to ensure the policy we are enforcing is the correct policy.
- Need to ensure we are enforcing policy according to appropriate principles under appropriate conditions.
 - Who are we showing the sensitive information to?
 - What computations have we done with the sensitive data before showing it?

Limitations: Password Example

from: HotCRP

to: Eve

subject: Password Reminder

Dear Eve,

Alice's password is [redacted].

<3, HotCRP

Limitations: Search Interface



Problem: Access Control Does Not Help Track Eventual *Viewer*

Ways viewer may be unpredictable:

- Viewer determined by user input.
 - “Send mail to...”
- Viewer computed from code.
 - Send to all users in a group.

Problem: Access Control Does Not Address *Implicit Flows*

```
int x := <secret>
if (x > 0) {
  y := y+1;
}
```

Information flow from x to y !

Ways Implicit Flows May Arise

What are some examples where we could capture information flows indirectly?

- Counting all users in a given location.
- Showing someone's photo in health record search results if some disease diagnosis is positive.

Goal: Track Sensitive Values



Lindsay Lohan with
parole ankle bracelet.

- Want to allow program to compute over sensitive values more or less freely.
- Want to prevent information from being released when there are unauthorized flows.


Some Questions

- What does access control give us?
- With access control, what is trusted?
- When isn't access control enough?
- What do we need to address the viewer problem and the problem of implicit flows?



Part Two: Process Isolation with OKWS (Krohn 2004)


Real-World Motivation: Online Dating Involves Secrets



Solution Requirements

- Needs to be able to run on Unix-based development server
- Needs to support all of the desired features of a production web server
- Needs to be fast enough to run OKCupid.com

Solution: Process Isolation



Run orthogonal services (for instance “search” and “inbox” in different processes.

This way, buffer overflow in “inbox” won’t affect “profile” or “search!”



Limitations

- Building secure systems with Unix is challenging because tools such as `setuid` and `chroot` conflict with common web server features such as embedded Python/Perl interpreters
- OKWS's security promises remain weak: if Bob comprises "inbox," can still read mail



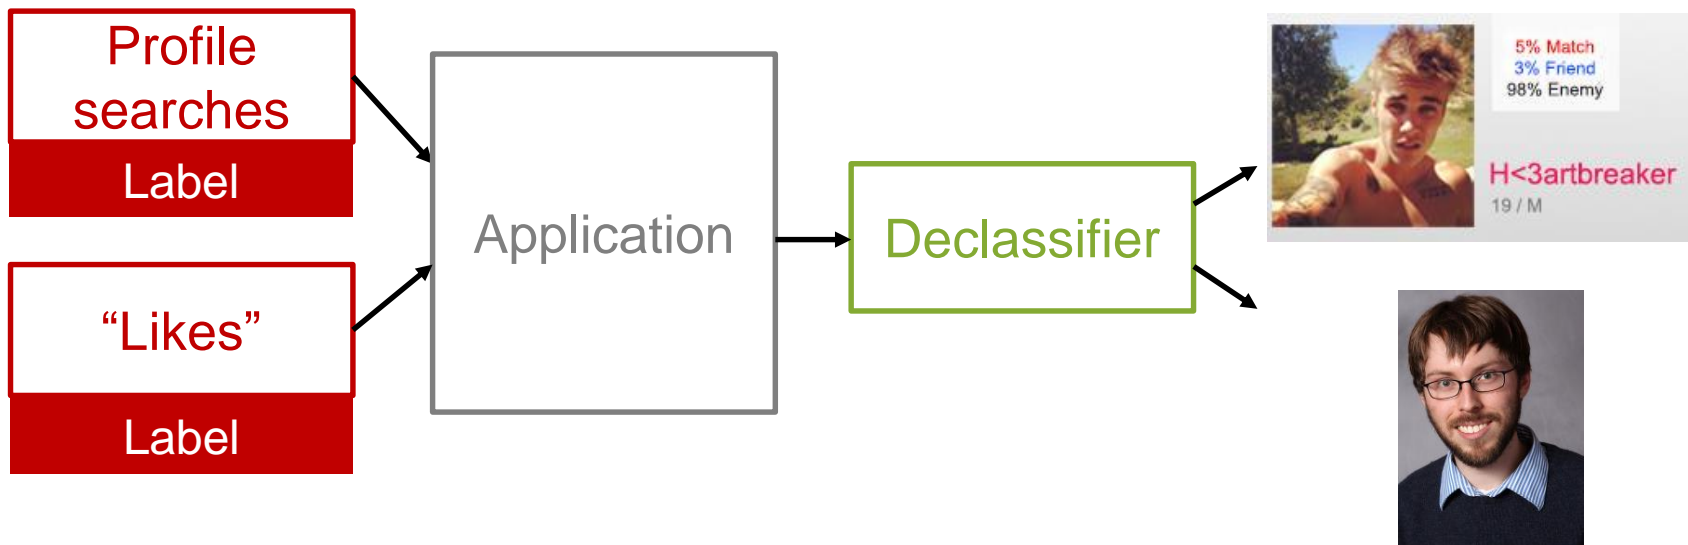
Part Three: Decentralized Information Flow Control with Flume (Krohn *et al* 2007)

What If We Could Run On a Customized OS?



Decentralized Information Flow


Model for controlling information flow in systems with *mutual distrust* and *decentralized authority*. Sensitive data is *labelled* and can be *declassified* in a decentralized way.



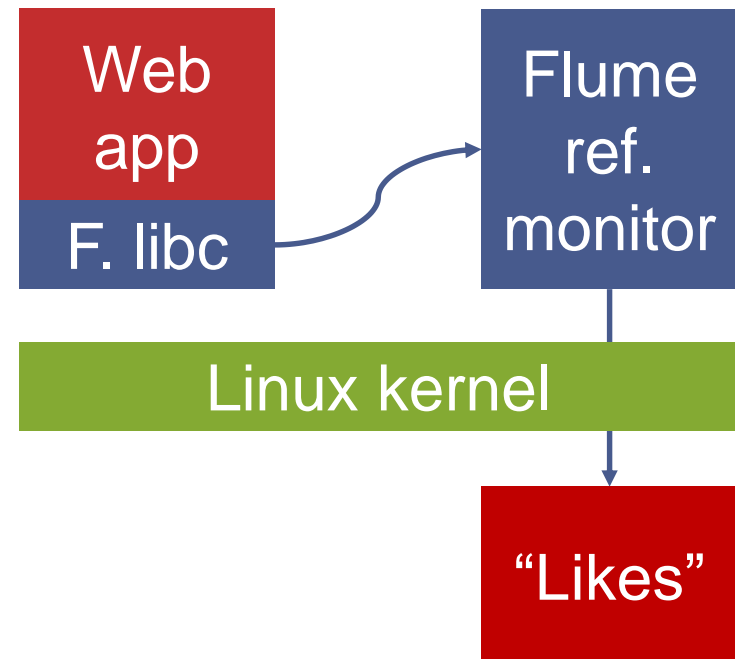
The Flume Operating System

Based on system call delegation. Built in user-space with a few small kernel patches on top of Linux and OpenBSD.

Regular OS

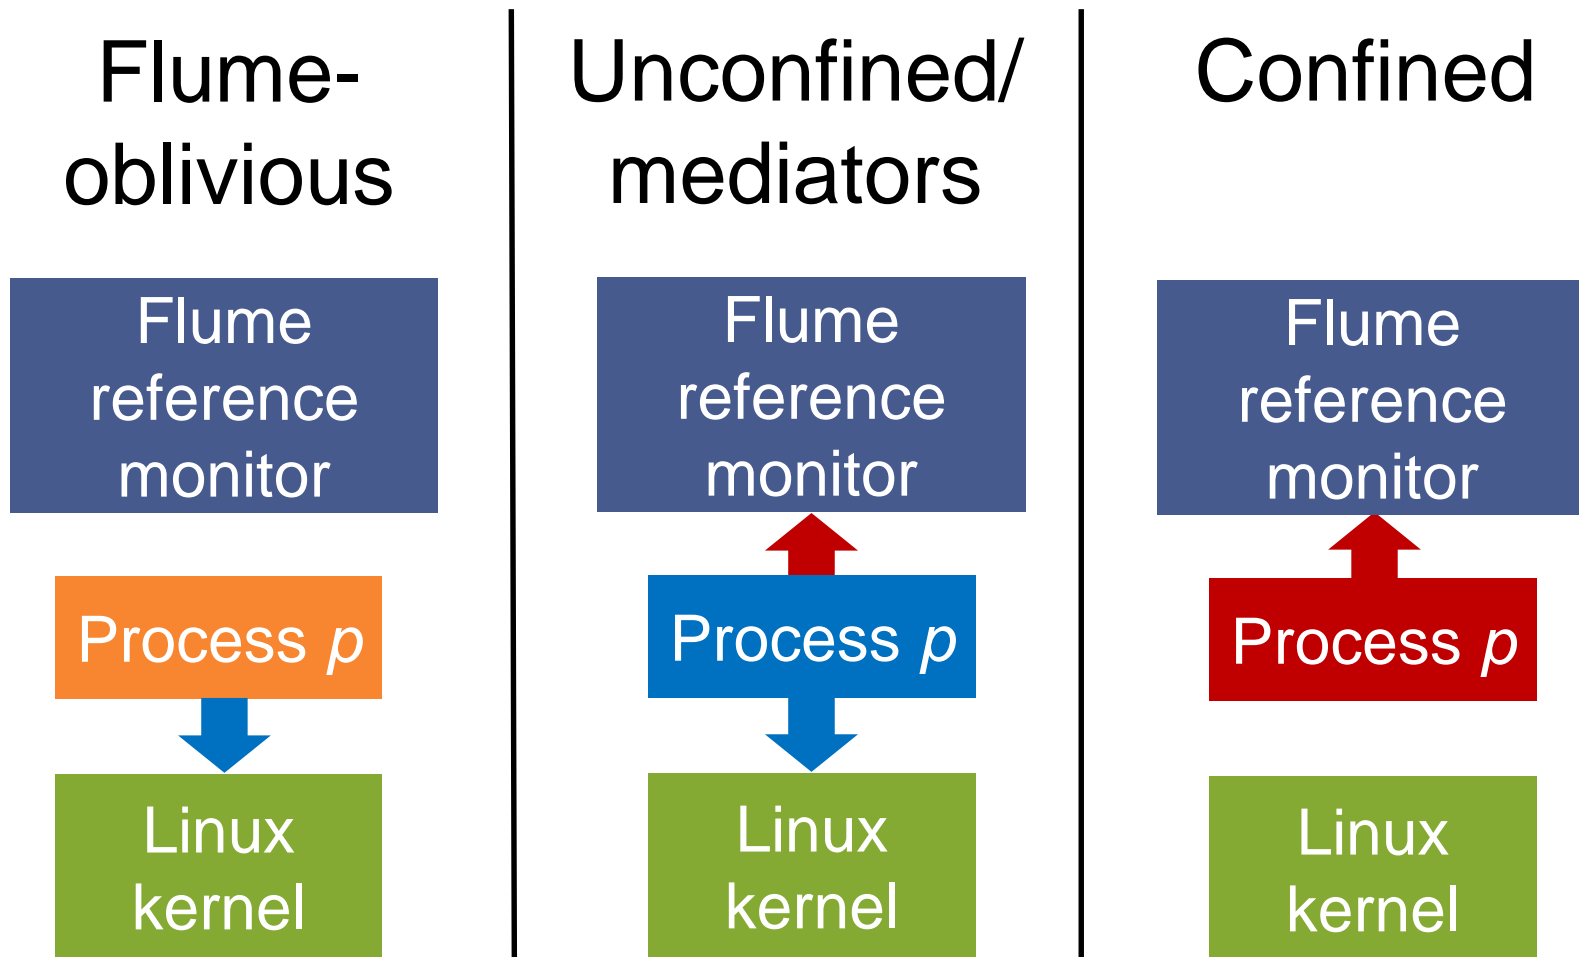


Flume OS



Three Classes of Processes

Based on slide by Max Krohn



Central Challenge

Accommodate process that use **existing communication interfaces** (for instance, socket and pipes) while specifying how and when they use their privileges.

- Awkward to modify each call to `read` or `write`.
- Conventional process interface full of channels that “leak” information, such as network sockets.

Solution: Labels + Endpoints

- *Label* processes with what they are allowed to read and write.
- Define how labels can be rewritten for *declassification* and *endorsement*.
- Represent each communication resource (for instance sockets and files) as an *endpoint* that specifies what subset of its privileges should be used when communicating.

Two Types of Processes

Untrusted

- Do most of the computation.
- Are constrained by, but possibly unaware of, DIFC controls.

Trusted

- Are aware of DIFC.
- Set up privacy and integrity controls that constrain untrusted processes.
- Have *privilege* to selectively violate classical information flow through *declassification* and *endorsement*.

Simple Label System

- **Goal:** track which secrets a process has accessed.
- **Mechanism:** each process gets a *secrecy label* summarizing the categories of data a process is assumed to have accessed.
 - { “Likes” } Tag
 - { “Financial reports” }
 - { “Likes” and “15-316 grades” } Label

Some Nomenclature

- **Confidentiality:** protecting sensitive reads.
 - When should a process be “authorized?”
 - Encryption provides end-to-end confidentiality, but it’s difficult to compute on encrypted data
- **Integrity:** protecting sensitive writes.
 - Only authorized processes can write a file
 - Digital signatures provide end-to-end integrity, but cannot change signed data

Confidentiality and Integrity

- Secrecy label (S_p)
 - Specifies what data process p has read
 - “/usr/bin/login may read the password file”
- Integrity label (I_p)
 - Used to **endorse** the trustworthiness of p
 - “/usr/bin/login can only be updated by root”
 - “/usr/bin/login can only read user libs and config files endorsed by root”

Privilege

Ownership (O_p) regulates how p can update S_p and I_p .

- **Endorsement:** tags p can add to its labels (e.g. t^+)
- **Declassification:** tags p can remove from its labels (e.g. t^-)
- D_p is the set of tags that p can both add and remove

Secrecy and Integrity, More Formally

- **Secrecy:** “At some point process p added data with tag s to its address space.”
 - $s \in S_p \Rightarrow \exists(data): p \text{ read data with tag } s$
- **Integrity:** “All inputs to process p had tag i .”
 - $i \in I_p \Rightarrow \forall(data): p \text{ read data with tag } i$

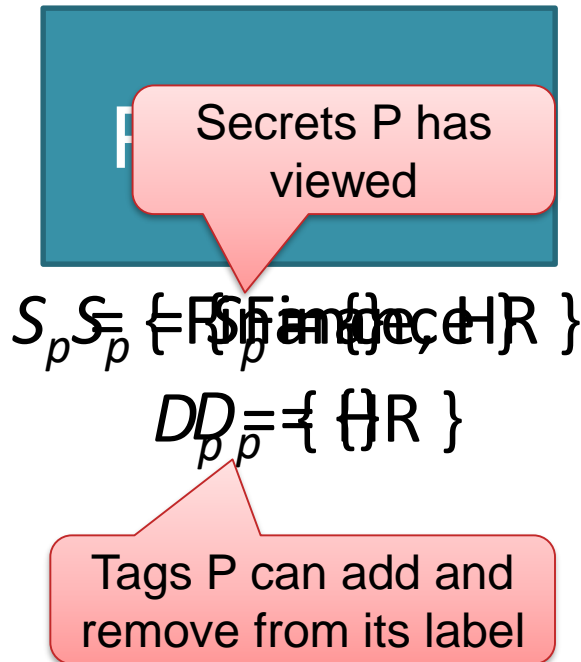
Privilege, More Formally

“ p can remove tag s from S_p and add tag i to I_p ”

- $s \in t^- \Rightarrow p$ is trusted to declassify s
- $i \in t^+ \Rightarrow p$ is trusted to endorse i
- $t \in D_p \Rightarrow t \in t^-$ and $t \in t^+$

Tags + Secrecy Labels

Based on slide by Max Krohn

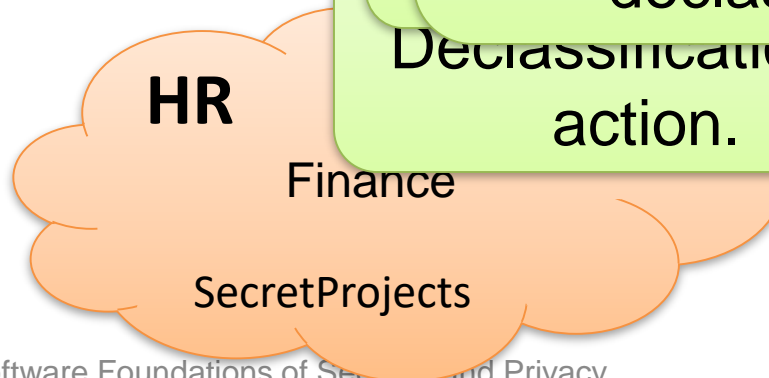


```
change_label({Finance});
tag_get_knowledge(tag);
change_label({R});
change_label({R});
```

DIFC Rule: A process can create a new tag; gets ability to declassify it.


Declassification in action.

Universe of Tags:



Tags + Integrity Labels

Based on slide by Max Krohn



$I_p = \{\text{Apple}\}$


$D_p = \{\}$

Tags P can add and remove from its label

`change_label1({});`

Any process can remove any tag from its label.

Universe of Tags:



Tags + Integrity Labels

Based on slide by Max Krohn


Process p

```
change_label1({});
```

$$I_p = \{\}$$

$$D_p = \{\}$$

Universe of Tags:



Tags + Integrity Labels


Based on slide by Max Krohn

Process p

$$I_p = \{\}$$
$$D_p = \{\}$$

```
change_label({});  
tag_get_base_integrity_tag(t);
```

Universe of Tags:



Tags + Integrity Labels

Based on slide by Max Krohn

Process p


```
change_label1({});
```

```
tag_t HR = create_tag();
```

$$I_p = \{\}$$

$$D_p = \{\}$$

Universe of Tags:



Tags + Integrity Labels

Based on slide by Max Krohn

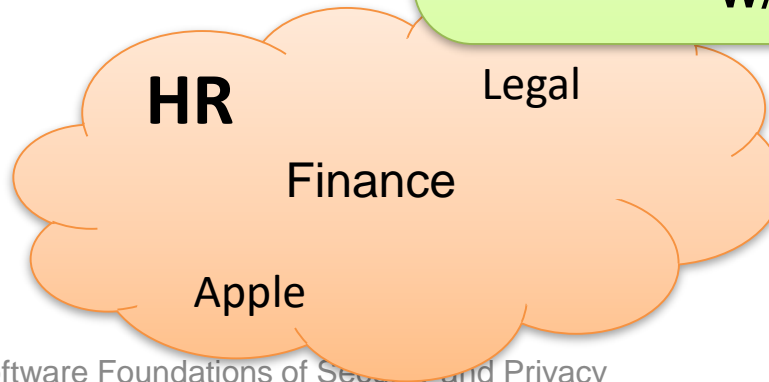
Process p

$$I_p = \{\}$$
$$D_p = \{HR\}$$

```
change_label1({});  
tag_t HR = create_tag();
```

DIFC Rule: A process can create a new tag; gets ability to endorse w/ it.

Universe of Tags:



Tags + Integrity Labels


Based on slide by Max Krohn

Process p

$$I_p = \{\}$$
$$D_p = \{\text{HR}\}$$

```
change_label1({});  
tag_t HR = create_tag();  
change_label1({HR});
```

Universe of Tags:



Tags + Integrity Labels

Based on slide by Max Krohn

Process p

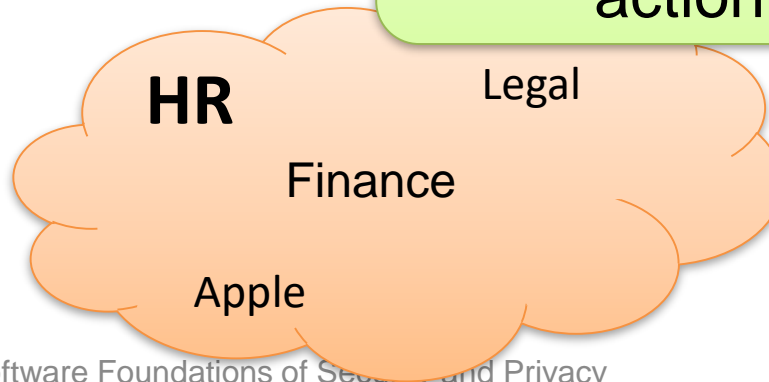
$I_p = \{HR\}$

$D_p = \{HR\}$

```
change_label({});  
tag_t HR = create_tag();  
change_label({HR});
```


DIFC:
Endorsement in
action.

Universe of Tags:



Communication Rule


Based on slide by Max Krohn



p can send to q iff $S_p \subseteq S_q$

Flume Communication Rule

Based on slide by Max Krohn

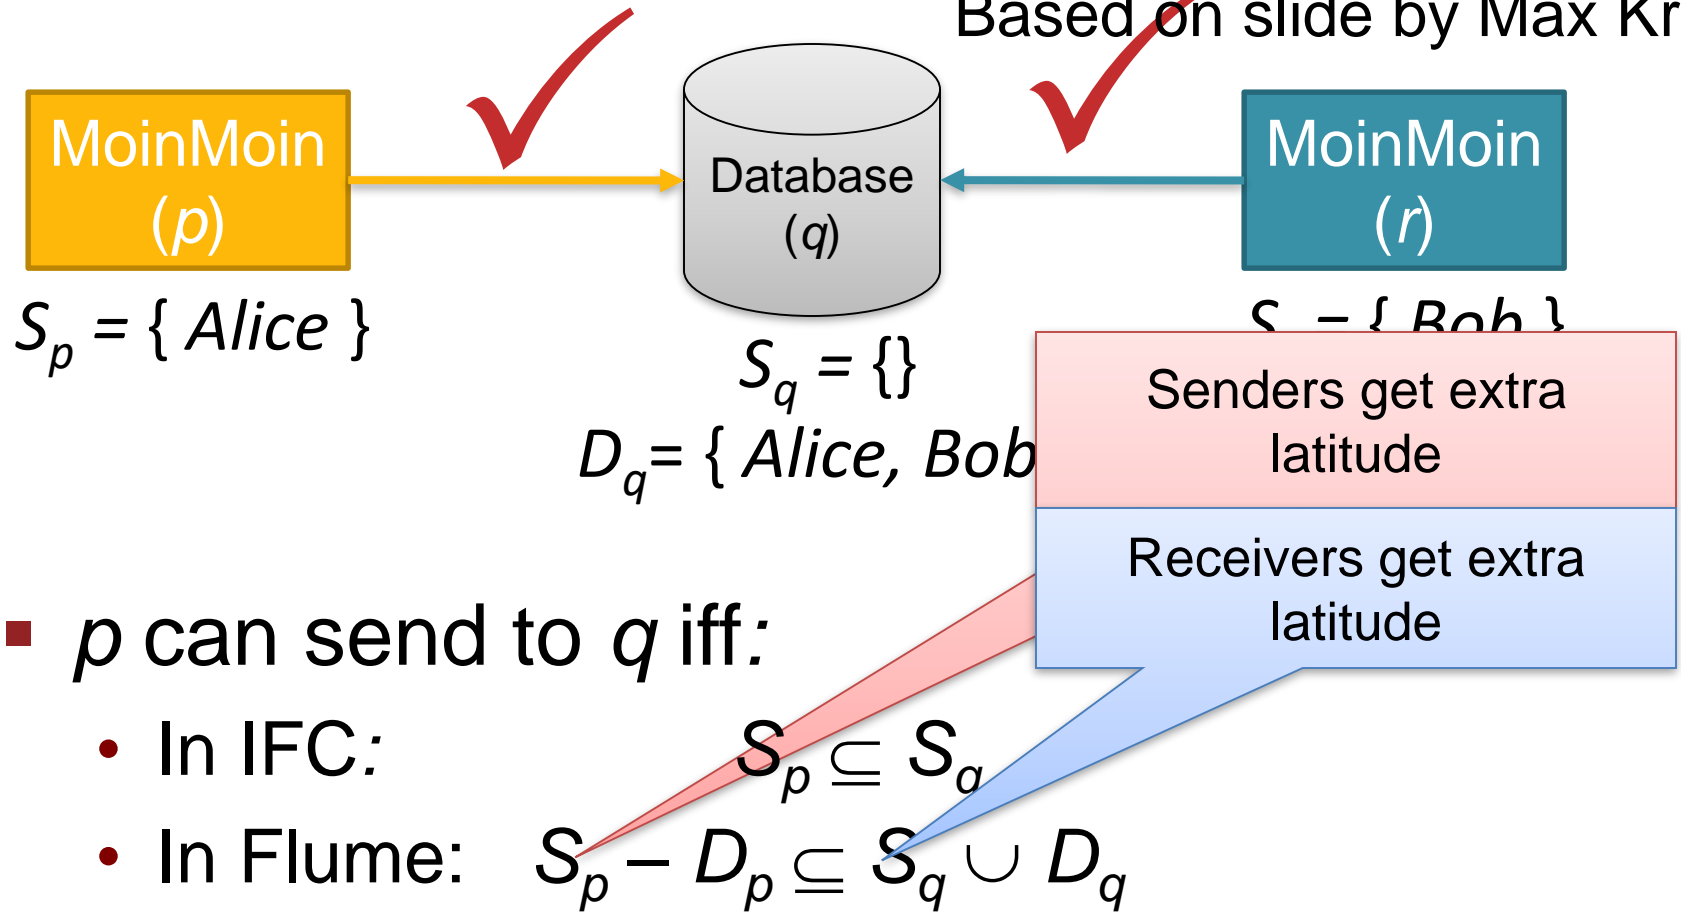


$$S_p \not\subseteq S_q$$

1. q changes to $S_q = \{ Alice \}$
2. p sends to q
3. q changes back to $S_q = \{ \}$

Flume Communication Rule

Based on slide by Max Krohn



The Unexpected Behavior Problem

Based on slide by Max Krohn




"Fire Alice, Bob, Charlie, Doug, Eddie, Frank, George, Hilda, Ilya..."



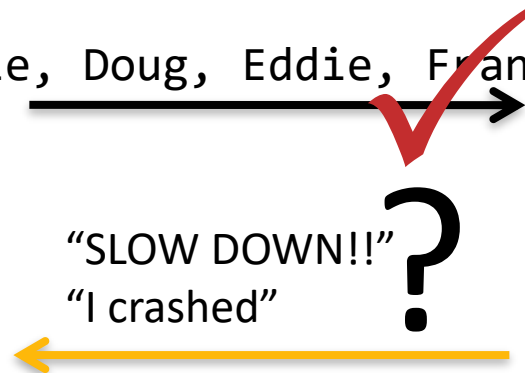
**Exposing failure messages
can leak information!**

The Unexpected Behavior Problem

Based on slide by Max Krohn




“Fire Alice, Bob, Charlie, Doug, Eddie, Frank, George, Hilda, Ilya...”



Solution: Endpoint Abstraction

Based on slide by Max Krohn




- “Fire Alice, Bob, Charlie, Doug, Eddie, Frank, George, Hilda, Ilya...”
- If $S_e \subseteq S_f$, then allow e to send to f
 - If $S_f \subseteq S_e$, then allow f to send to e
 - If $S_f = S_e$, then allow bidirectional flow
- “SLOW DOWN!!”
- “I crashed”

Benefits of Endpoints

- Simplifies application programming. When a process attempts and fails to adjust labels on its endpoints, system can safely report errors.
- Make many declassification and endorsement decisions *explicit*. Flume processes must explicitly mark file descriptors that serve as avenues for declassification/endorsement.

Endpoints Declassify Data

Based on slide by Max Krohn



Data enters
process p with
secrecy $\{ \text{HR} \}$

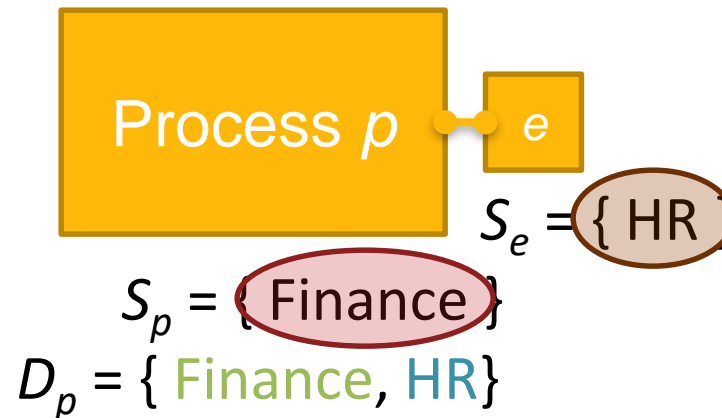
But p keeps its
label $S_p = \{ \}$

Thus p needs
 $\text{HR} \in D_p$

Endpoint Invariant

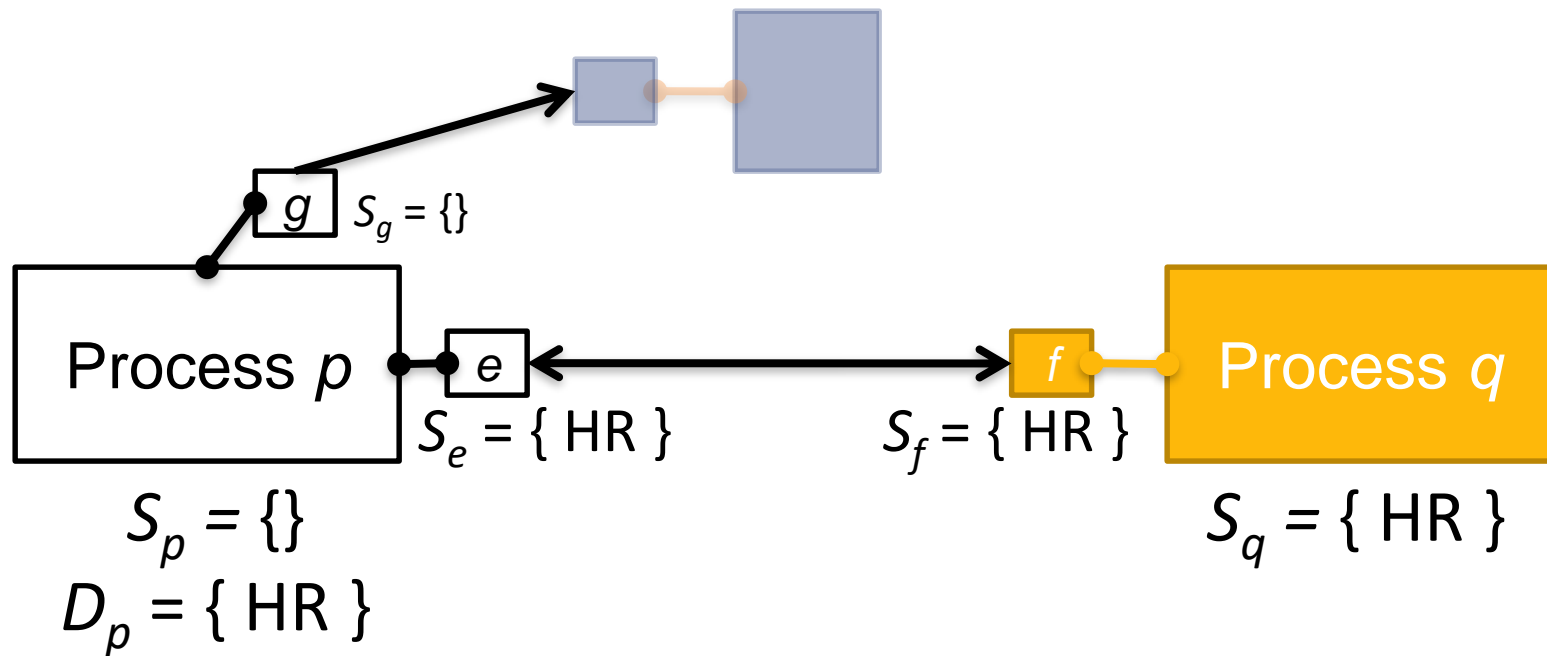
Based on slides by Frank Krohn

- For any tag $t \in S_p$ and $t \notin S_e$ Export inf.
- Or any tag $t \in S_e$ and $t \notin S_p$ Import inf.
- It must be that $t \in D_p$



Endpoints Labels Are Independent

Based on slide by Max Krohn



Evaluation

Based on slide by Max Krohn

- Does Flume allow adoption of Unix software?
 - 1,000 LOC launcher/declassifier
 - 1,000 out of 100,000 LOC in MoinMoin changed
 - Python interpreter, Apache, unchanged
- Does Flume solve security vulnerabilities?
 - Without our knowing, MoinMoin wiki case studies inherited two ACL bypass bugs from MoinMoin
 - Both are not exploitable in Flume's MoinMoin
- Does Flume perform reasonably?
 - Performs within a factor of 2 of the original on read and write benchmarks

Limitations

Based on slide by Max Krohn

- Bigger TCB than HiStar / Asbestos
 - Linux stack (Kernel + glibc + linker)
 - Reference monitor (~22 kLOC)
- Covert channels via disk quotas
- Confined processes like MoinMoin don't get full POSIX API.
 - `spawn()` instead of `fork()` & `exec()`
 - `flume_pipe()` instead of `pipe()`



Part Four: How Do We Have a Reference Monitor for IFC?

Recall: Information Flow Isn't EM-Enforceable

Let S_1 :

```
if (x)
  y = 0;
else
  y = 1;
```

And S_2 :

```
x, y = 0, 1;
```

And S_3 :

```
x, y = 1, 0;
```

- Recall that we can only distinguish between these three programs if we can choose *two* values for x .
- Information flow is a *hyperproperty*.
- Information flow is a *2-safety* property that is finitely refutable over *pairs* of traces.

Desired Guarantee

Non-interference: observable program behavior should not depend on confidential data.

More formally, for a deterministic program P :

$$\begin{aligned} \forall M_1, M_2: \quad & M_1 =_L M_2 \wedge && \text{Low-confidentiality} \\ & (P, M_1) \rightarrow^* M'_1 \wedge && \text{projections of initial memory} \\ & (P, M_2) \rightarrow^* M'_2 \Rightarrow && \text{are equivalent.} \\ & M'_1 =_L M'_2 && \text{Low-confidentiality} \\ & && \text{projections of result memory} \\ & && \text{are equivalent.} \end{aligned}$$

But Also Recall!

```
while(read(&buf, &len, fp)) {  
    if(buf[0] == 255)  
        send(sock, buf, len);  
    printf("%s", buf);  
}
```

```
while(read(&buf, &len, fp)) {  
    memset(buf, 0, len);  
    send(sock, buf, len);  
    printf("%s", buf);  
}
```

Does this flow fp to sock?

- We discussed “no send after read” policy earlier.
- Ideally want to prevent fp from flowing into sock.
- But second program does not flow fp into sock! Check was conservative.

“Platonic” Information Flow is Fine-Grained!

Things that should be allowed under our definition of non-interference:

- Sensitive value gets sent to process, but it doesn't actually use it.
- Sensitive value gets sent to process and it uses it, but doesn't send it out across a specific endpoint.

So What Is Flume Doing?

- Flume **overapproximates** programs that can leak information.
- Flume tracks information flow at the **process level**, rather than at the granularity of individual reads/writes.

Question: What flows does Flume prevent, that may be otherwise allowed?



Part Five: Wrapping Up on Coarse-Grained IFC

Information Flow Involves Tracking Across the Program



Internet joke about how Lindsay Lohan tried to cover up her bracelet.

In a DIFC system, it is not only sensitive values that are tracked, but *any* value whose value depends on a sensitive value. Tracking becomes very fashionable!

Granularity Matters

- Precise information flow is not a safety property.
- We can, however, *overapproximate* information flow using reference monitors.
- There are tradeoffs between precision of tracking and programmer/runtime overhead!

Discussion Questions

- How does access control fall short? (Why do we need information flow?)
- What are the tradeoffs of the label abstraction? The endpoint abstraction?
- What are the tradeoffs of using DFIC systems in general?
- How do the techniques we learned about prevent the leaks we discussed?